

Sockets

Sockets provide two-way *interprocess communications* between processes on one or more machines.

Sockets

Sockets provide two-way *interprocess communications* between processes on one or more machines.

connectionless each packet contains full addresses. (*datagrams*)

Sockets

Sockets provide two-way *interprocess communications* between processes on one or more machines.

connectionless each packet contains full addresses. (*datagrams*)

connection-based sender/receiver (*ie. server/client*) addresses used to establish a connection which simplifies (*ie. shortens*) packets' headers; a **virtual circuit** is established.

(data appear to be transmitted in an unbroken stream; reliable)

Sockets

Sockets provide two-way *interprocess communications* between processes on one or more machines.

connectionless each packet contains full addresses. (*datagrams*)

connection-based sender/receiver (*ie. server/client*) addresses used to establish a connection which simplifies (*ie. shortens*) packets' headers; a **virtual circuit** is established.

(data appear to be transmitted in an unbroken stream; reliable)

Sockets of both types must specify a **5-tuple**:

Sockets

Sockets provide two-way *interprocess communications* between processes on one or more machines.

connectionless each packet contains full addresses. (*datagrams*)

connection-based sender/receiver (*ie. server/client*) addresses used to establish a connection which simplifies (*ie. shortens*) packets' headers; a **virtual circuit** is established.

(data appear to be transmitted in an unbroken stream; reliable)

Sockets of both types must specify a **5-tuple**:

| | Type | Protocol | Local-Addr | Local-Process | Foreign-Addr | Foreign-Process |
|--|----------------------------|----------|------------|---------------|--------------------|-----------------|
| | connection oriented server | socket() | bind() | | listen(), accept() | |
| | connection oriented client | socket() | connect() | | | |
| | connectionless server | socket() | bind() | | recvfrom() | |
| | connectionless client | socket() | bind() | | sendto() | |

Connection-Oriented Setup

Server

socket()

Client

Description

creates socket: family, domain, protocol (*1-tuple*)

Connection-Oriented Setup

Server

Client

Description

socket()

creates socket: family, domain, protocol (*1-tuple*)

bind()

fills in local addr, local process (*3-tuple*)

Connection-Oriented Setup

Server

socket()

bind()

listen()

Client

Description

creates socket: family, domain, protocol (*1-tuple*)

fills in local addr, local process (*3-tuple*)

server specifies “calls waiting”:

ie. qty clients that may queue for connection

Connection-Oriented Setup

Server

Client

Description

socket()

creates socket: family, domain, protocol (*1-tuple*)

bind()

fills in local addr, local process (*3-tuple*)

listen()

server specifies “calls waiting”:

ie. qty clients that may queue for connection

accept()

creates a new 5-tuple with foreign-addr,process

blocks

until

connect

Connection-Oriented Setup

Server

socket()

bind()

listen()

accept()

blocks

until

connect

Client

socket()

Description

creates socket: family, domain, protocol (*1-tuple*)

fills in local addr, local process (*3-tuple*)

server specifies “calls waiting”:

ie. qty clients that may queue for connection

creates a new 5-tuple with foreign-addr,process

creates socket: family, domain, protocol (*1-tuple*)

Connection-Oriented Setup

Server

socket()

bind()

listen()

accept()

blocks

until

connect

Client

socket()

connect()

Description

creates socket: family, domain, protocol (*1-tuple*)

fills in local addr, local process (*3-tuple*)

server specifies “calls waiting”:

ie. qty clients that may queue for connection

creates a new 5-tuple with foreign-addr,process

creates socket: family, domain, protocol (*1-tuple*)

establishes connection to server (*5-tuple!*)

Connection-Oriented Setup

Server

socket()

bind()

listen()

accept()

blocks

until

connect

write()/read()
send()/recv()

Client

socket()

connect()

write()/read()
send()/recv()



Description

creates socket: family, domain, protocol (*1-tuple*)

fills in local addr, local process (*3-tuple*)

server specifies “calls waiting”:

ie. qty clients that may queue for connection

creates a new 5-tuple with foreign-addr,process

creates socket: family, domain, protocol (*1-tuple*)

establishes connection to server (*5-tuple!*)

communications

Connection-Oriented Setup

Server

socket()

bind()

listen()

accept()

blocks

until

connect

write()/read()

send()/recv()

shutdown()

Client

socket()

connect()

write()/read()

send()/recv()

shutdown()



Description

creates socket: family, domain, protocol (*1-tuple*)

fills in local addr, local process (*3-tuple*)

server specifies “calls waiting”:

ie. qty clients that may queue for connection

creates a new 5-tuple with foreign-addr,process

creates socket: family, domain, protocol (*1-tuple*)

establishes connection to server (*5-tuple!*)

communications

shuts down all or part of connection

Connection-Oriented Setup

Server

socket()

bind()

listen()

accept()

blocks

until

connect

write()/read()
send()/recv()

shutdown()

close()

Client

socket()

connect()

write()/read()
send()/recv()

shutdown()

close()



Description

creates socket: family, domain, protocol (*1-tuple*)

fills in local addr, local process (*3-tuple*)

server specifies “calls waiting”:

ie. qty clients that may queue for connection

creates a new 5-tuple with foreign-addr,process

creates socket: family, domain, protocol (*1-tuple*)

establishes connection to server (*5-tuple!*)

communications

shuts down all or part of connection

closes associated file descriptor

Connectionless Setup

socket()

socket()

specifies protocol (*1-tuple*)

Connectionless Setup

socket()

socket()

specifies protocol (*1-tuple*)

bind()

bind()

specifies local addr,process (*3-tuple*)

Connectionless Setup

socket()

socket()

specifies protocol (*1-tuple*)

bind()

bind()

specifies local addr,process (*3-tuple*)

recvfrom()/sendto()

⇔

recvfrom()/sendto()

specifies (*5-tuple*)

foreign addr,process

Connectionless Setup

socket()

socket()

specifies protocol (*1-tuple*)

bind()

bind()

specifies local addr,process (*3-tuple*)

recvfrom()/sendto()

\iff

recvfrom()/sendto()

specifies (*5-tuple*)

foreign addr,process

- Connectionless packets may get routed different ways and so may arrive out of sequence

Connectionless Setup

socket()

socket()

specifies protocol (*1-tuple*)

bind()

bind()

specifies local addr,process (*3-tuple*)

recvfrom()/sendto()

⇔

recvfrom()/sendto()

specifies (*5-tuple*)

foreign addr,process

- Connectionless packets may get routed different ways and so may arrive out of sequence
- Packets compete for network bandwidth and may be dropped due to “collisions”

Connectionless Setup

socket()

socket()

specifies protocol (*1-tuple*)

bind()

bind()

specifies local addr,process (*3-tuple*)

recvfrom()/sendto()

⇔

recvfrom()/sendto()

specifies (*5-tuple*)

foreign addr,process

- Connectionless packets may get routed different ways and so may arrive out of sequence
- Packets compete for network bandwidth and may be dropped due to “collisions”
- Connection-oriented communications are built from connectionless services

Connectionless Setup

socket()

socket()

specifies protocol (*1-tuple*)

bind()

bind()

specifies local addr,process (*3-tuple*)

recvfrom()/sendto()

⇔

recvfrom()/sendto()

specifies (*5-tuple*)

foreign addr,process

- Connectionless packets may get routed different ways and so may arrive out of sequence
- Packets compete for network bandwidth and may be dropped due to “collisions”
- Connection-oriented communications are built from connectionless services
- Considerable effort has gone to eliminate these issues from connection oriented packets

Connectionless Setup

socket()

socket()

specifies protocol (*1-tuple*)

bind()

bind()

specifies local addr,process (*3-tuple*)

recvfrom()/sendto()

⇔

recvfrom()/sendto()

specifies (*5-tuple*)

foreign addr,process

- Connectionless packets may get routed different ways and so may arrive out of sequence
- Packets compete for network bandwidth and may be dropped due to “collisions”
- Connection-oriented communications are built from connectionless services
- Considerable effort has gone to eliminate these issues from connection oriented packets
- If you want reliable, properly sequenced packets, don't re-invent the wheel: use connection-oriented communications

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

| | | |
|---------|---------|---------------|
| domains | AF_UNIX | unix internal |
|---------|---------|---------------|

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

| | | |
|---------|---------|---------------|
| domains | AF_UNIX | unix internal |
| | AF_INET | internet |

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

| | | |
|---------|---------|--------------------------|
| domains | AF_UNIX | unix internal |
| | AF_INET | internet |
| | AF_NS | Novell/Xerox NS Protocol |

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

| | | |
|---------|------------|--------------------------|
| domains | AF_UNIX | unix internal |
| | AF_INET | internet |
| | AF_NS | Novell/Xerox NS Protocol |
| | AF_IMPLINK | IMP link layer |

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

| | | |
|---------|-------------|------------------------------|
| domains | AF_UNIX | unix internal |
| | AF_INET | internet |
| | AF_NS | Novell/Xerox NS Protocol |
| | AF_IMPLINK | IMP link layer |
| types | SOCK_STREAM | stream (connection oriented) |

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

| | | |
|---------|-------------|------------------------------|
| domains | AF_UNIX | unix internal |
| | AF_INET | internet |
| | AF_NS | Novell/Xerox NS Protocol |
| | AF_IMPLINK | IMP link layer |
| types | SOCK_STREAM | stream (connection oriented) |
| | SOCK_DGRAM | datagram (connectionless) |

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

| | | |
|---------|-------------|------------------------------|
| domains | AF_UNIX | unix internal |
| | AF_INET | internet |
| | AF_NS | Novell/Xerox NS Protocol |
| | AF_IMPLINK | IMP link layer |
| types | SOCK_STREAM | stream (connection oriented) |
| | SOCK_DGRAM | datagram (connectionless) |
| | SOCK_RAW | raw network protocol |

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

| | | |
|---------|----------------|---|
| domains | AF_UNIX | unix internal |
| | AF_INET | internet |
| | AF_NS | Novell/Xerox NS Protocol |
| | AF_IMPLINK | IMP link layer |
| types | SOCK_STREAM | stream (connection oriented) |
| | SOCK_DGRAM | datagram (connectionless) |
| | SOCK_RAW | raw network protocol |
| | SOCK_SEQPACKET | sequenced reliable 2-way packet, fixed max length |

socket()

```
#include <sys/types.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

(1-tuple)

- Returns -1 on failure, + for socket descriptor
- “AF” stands for “*Address Family*”:

| | | |
|---------|----------------|---|
| domains | AF_UNIX | unix internal |
| | AF_INET | internet |
| | AF_NS | Novell/Xerox NS Protocol |
| | AF_IMPLINK | IMP link layer |
| types | SOCK_STREAM | stream (connection oriented) |
| | SOCK_DGRAM | datagram (connectionless) |
| | SOCK_RAW | raw network protocol |
| | SOCK_SEQPACKET | sequenced reliable 2-way packet, fixed max length |
| | SOCK_RDM | reliably delivered packet, (but not sequencing) |

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

Socket types are not implemented by all protocol families:

Valid Type/Domain/Protocol Combinations

| <i>type</i> | <i>AF_UNIX</i> | <i>AF_INET</i> | <i>AF_NS</i> |
|-------------|----------------|----------------|--------------|
| | | | |

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

Socket types are not implemented by all protocol families:

Valid Type/Domain/Protocol Combinations

| <i>type</i> | <i>AF_UNIX</i> | <i>AF_INET</i> | <i>AF_NS</i> |
|----------------|----------------|----------------|--------------|
| SOCK_STREAM | | | |
| SOCK_DGRAM | | | |
| SOCK_RAW | | | |
| SOCK_SEQPACKET | | | |

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

Socket types are not implemented by all protocol families:

Valid Type/Domain/Protocol Combinations

| <i>type</i> | <i>AF_UNIX</i> | <i>AF_INET</i> | <i>AF_NS</i> |
|----------------|----------------|----------------|--------------|
| SOCK_STREAM | ✓ | | |
| SOCK_DGRAM | ✓ | | |
| SOCK_RAW | | | |
| SOCK_SEQPACKET | | | |

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

Socket types are not implemented by all protocol families:

Valid Type/Domain/Protocol Combinations

| <i>type</i> | <i>AF_UNIX</i> | <i>AF_INET</i> | <i>AF_NS</i> |
|----------------|----------------|----------------|--------------|
| SOCK_STREAM | ✓ | TCP | |
| SOCK_DGRAM | ✓ | UDP | |
| SOCK_RAW | | IP | |
| SOCK_SEQPACKET | | | |

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

Socket types are not implemented by all protocol families:

Valid Type/Domain/Protocol Combinations

| <i>type</i> | <i>AF_UNIX</i> | <i>AF_INET</i> | <i>AF_NS</i> |
|----------------|----------------|----------------|--------------|
| SOCK_STREAM | ✓ | TCP | SPP |
| SOCK_DGRAM | ✓ | UDP | IDP |
| SOCK_RAW | | IP | ✓ |
| SOCK_SEQPACKET | | | SPP |

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

Socket types are not implemented by all protocol families:

Valid Type/Domain/Protocol Combinations

| <i>type</i> | <i>AF_UNIX</i> | <i>AF_INET</i> | <i>AF_NS</i> |
|----------------|----------------|----------------|--------------|
| SOCK_STREAM | ✓ | TCP | SPP |
| SOCK_DGRAM | ✓ | UDP | IDP |
| SOCK_RAW | | IP | ✓ |
| SOCK_SEQPACKET | | | SPP |

- Sockets of type SOCK_STREAM are full-duplex byte streams and resemble pipes.

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

Socket types are not implemented by all protocol families:

Valid Type/Domain/Protocol Combinations

| <i>type</i> | <i>AF_UNIX</i> | <i>AF_INET</i> | <i>AF_NS</i> |
|----------------|----------------|----------------|--------------|
| SOCK_STREAM | ✓ | TCP | SPP |
| SOCK_DGRAM | ✓ | UDP | IDP |
| SOCK_RAW | | IP | ✓ |
| SOCK_SEQPACKET | | | SPP |

- Sockets of type SOCK_STREAM are full-duplex byte streams and resemble pipes.
- SOCK_STREAM communications do not preserve record boundaries

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

Socket types are not implemented by all protocol families:

Valid Type/Domain/Protocol Combinations

| <i>type</i> | <i>AF_UNIX</i> | <i>AF_INET</i> | <i>AF_NS</i> |
|----------------|----------------|----------------|--------------|
| SOCK_STREAM | ✓ | TCP | SPP |
| SOCK_DGRAM | ✓ | UDP | IDP |
| SOCK_RAW | | IP | ✓ |
| SOCK_SEQPACKET | | | SPP |

- Sockets of type SOCK_STREAM are full-duplex byte streams and resemble pipes.
- SOCK_STREAM communications do not preserve record boundaries
- SOCK_DGRAMs do preserve record boundaries; they are not reliable nor sequenced

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

Socket types are not implemented by all protocol families:

Valid Type/Domain/Protocol Combinations

| <i>type</i> | <i>AF_UNIX</i> | <i>AF_INET</i> | <i>AF_NS</i> |
|----------------|----------------|----------------|--------------|
| SOCK_STREAM | ✓ | TCP | SPP |
| SOCK_DGRAM | ✓ | UDP | IDP |
| SOCK_RAW | | IP | ✓ |
| SOCK_SEQPACKET | | | SPP |

- Sockets of type SOCK_STREAM are full-duplex byte streams and resemble pipes.
- SOCK_STREAM communications do not preserve record boundaries
- SOCK_DGRAMs do preserve record boundaries; they are not reliable nor sequenced
- SOCK_RAW sockets resemble SOCK_DGRAMs; they send datagrams via sendto()/recvfrom() (common use: ping)

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

protocol: usually zero (socket() will pick an appropriate one)

| <i>Family</i> | <i>Type</i> | <i>Protocol</i> | <i>Actual Protocol</i> |
|---------------|-------------|-----------------|------------------------|
| | | | |
| | | | |

socket(), con't.

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol)
```

protocol: usually zero (socket() will pick an appropriate one)

| <i>Family</i> | <i>Type</i> | <i>Protocol</i> | <i>Actual Protocol</i> |
|---------------|-------------|-----------------|------------------------|
| AF_INET | | | |
| AF_NS | | | |

socket(), con't.

#include <sys/types.h>

#include <sys/socket.h>

int socket(int domain, int type, int protocol)

protocol: usually zero (socket() will pick an appropriate one)

| <i>Family</i> | <i>Type</i> | <i>Protocol</i> | <i>Actual Protocol</i> |
|---------------|----------------|-----------------|------------------------|
| AF_INET | SOCK_DGRAM | | |
| | SOCK_STREAM | | |
| | SOCK_RAW | | |
| | SOCK_RAW | | |
| AF_NS | SOCK_STREAM | | |
| | SOCK_SEQPACKET | | |
| | SOCK_RAW | | |
| | SOCK_RAW | | |

socket(), con't.

#include <sys/types.h>

#include <sys/socket.h>

int socket(int domain, int type, int protocol)

protocol: usually zero (socket() will pick an appropriate one)

| <i>Family</i> | <i>Type</i> | <i>Protocol</i> | <i>Actual Protocol</i> |
|---------------|----------------|-----------------|------------------------|
| AF_INET | SOCK_DGRAM | IPPROTO_UDP | |
| | SOCK_STREAM | IPPROTO_TCP | |
| | SOCK_RAW | IPPROTO_ICMP | |
| | SOCK_RAW | IPPROTO_RAW | |
| AF_NS | SOCK_STREAM | NSPROTO_SPP | |
| | SOCK_SEQPACKET | NSPROTO_SPP | |
| | SOCK_RAW | NSPROTO_ERROR | |
| | SOCK_RAW | NSPROTO_RAW | |

socket(), con't.

#include <sys/types.h>

#include <sys/socket.h>

int socket(int domain, int type, int protocol)

protocol: usually zero (socket() will pick an appropriate one)

| <i>Family</i> | <i>Type</i> | <i>Protocol</i> | <i>Actual Protocol</i> |
|---------------|----------------|-----------------|------------------------|
| AF_INET | SOCK_DGRAM | IPPROTO_UDP | UDP |
| | SOCK_STREAM | IPPROTO_TCP | TCP |
| | SOCK_RAW | IPPROTO_ICMP | ICMP |
| | SOCK_RAW | IPPROTO_RAW | (raw) |
| AF_NS | SOCK_STREAM | NSPROTO_SPP | SPP |
| | SOCK_SEQPACKET | NSPROTO_SPP | SPP |
| | SOCK_RAW | NSPROTO_ERROR | (error) |
| | SOCK_RAW | NSPROTO_RAW | (raw) |

Socket Definitions

UDP

User Data Protocol (connectionless, no guarantee of delivery, uses IP)

Socket Definitions

- UDP User Data Protocol (connectionless, no guarantee of delivery, uses IP)
- TCP Transmission Control Protocol (connection, reliable, full duplex, byte stream, uses IP)

Socket Definitions

- UDP User Data Protocol (connectionless, no guarantee of delivery, uses IP)
- TCP Transmission Control Protocol (connection, reliable, full duplex, byte stream, uses IP)
- ICMP Internet Control Message Protocol (handles error and control between gateways and hosts, uses IP)

Socket Definitions

| | |
|------|---|
| UDP | User Data Protocol (connectionless, no guarantee of delivery, uses IP) |
| TCP | Transmission Control Protocol (connection, reliable, full duplex, byte stream, uses IP) |
| ICMP | Internet Control Message Protocol (handles error and control between gateways and hosts, uses IP) |
| IP | Internet Protocol. Provides packet delivery for TCP, UDP, ICMP. |

Socket Definitions

| | |
|------|---|
| UDP | User Data Protocol (connectionless, no guarantee of delivery, uses IP) |
| TCP | Transmission Control Protocol (connection, reliable, full duplex, byte stream, uses IP) |
| ICMP | Internet Control Message Protocol (handles error and control between gateways and hosts, uses IP) |
| IP | Internet Protocol. Provides packet delivery for TCP, UDP, ICMP. |
| ARP | Address Resolution Protocol (maps internet address to hardware address) |

Socket Definitions

| | |
|------|---|
| UDP | User Data Protocol (connectionless, no guarantee of delivery, uses IP) |
| TCP | Transmission Control Protocol (connection, reliable, full duplex, byte stream, uses IP) |
| ICMP | Internet Control Message Protocol (handles error and control between gateways and hosts, uses IP) |
| IP | Internet Protocol. Provides packet delivery for TCP, UDP, ICMP. |
| ARP | Address Resolution Protocol (maps internet address to hardware address) |
| RARP | Reverse Address Resolution Protocol (maps hardware address to internet address) |

Socket Definitions

| | |
|------|---|
| UDP | User Data Protocol (connectionless, no guarantee of delivery, uses IP) |
| TCP | Transmission Control Protocol (connection, reliable, full duplex, byte stream, uses IP) |
| ICMP | Internet Control Message Protocol (handles error and control between gateways and hosts, uses IP) |
| IP | Internet Protocol. Provides packet delivery for TCP, UDP, ICMP. |
| ARP | Address Resolution Protocol (maps internet address to hardware address) |
| RARP | Reverse Address Resolution Protocol (maps hardware address to internet address) |
| SPP | Sequenced Packet Protocol (a Xerox Network Systems protocol for reliable flow-controlled packet delivery) |

Open Systems Interconnection Model

- The International Organization for Standardization (*ISO*) developed an *abstract model* of networking (called the *Basic Reference Model*)

Open Systems Interconnection Model

- The International Organization for Standardization (*ISO*) developed an *abstract model* of networking (called the *Basic Reference Model*)
- The model consists of seven layers (*see next slide*)

Open Systems Interconnection Model

- The International Organization for Standardization (*ISO*) developed an *abstract model* of networking (called the *Basic Reference Model*)
- The model consists of seven layers (*see next slide*)
- Each layer only interacts directly with the layer directly beneath it

OSI Layers

Open Systems Interconnection Model

| <i>Data Unit</i> | <i>Layer</i> | <i>Function</i> |
|------------------|----------------|--------------------------------|
| Data | 7. Application | Network process to application |

OSI Layers

Open Systems Interconnection Model

| <i>Data Unit</i> | <i>Layer</i> | <i>Function</i> |
|------------------|-----------------|--|
| Data | 7. Application | Network process to application |
| Data | 6. Presentation | Data representation, encryption/decryption, convert dependent data to machine independent data |

OSI Layers

Open Systems Interconnection Model

| <i>Data Unit</i> | <i>Layer</i> | <i>Function</i> |
|------------------|-----------------|--|
| Data | 7. Application | Network process to application |
| Data | 6. Presentation | Data representation, encryption/decryption, convert dependent data to machine independent data |
| Data | 5. Session | Interhost communication, managing sessions between applications |

OSI Layers

Open Systems Interconnection Model

| <i>Data Unit</i> | <i>Layer</i> | <i>Function</i> |
|------------------|-----------------|--|
| Data | 7. Application | Network process to application |
| Data | 6. Presentation | Data representation, encryption/decryption, convert dependent data to machine independent data |
| Data | 5. Session | Interhost communication, managing sessions between applications |
| Segments | 4. Transport | End-to-end connections, reliability, and flow control |

OSI Layers

Open Systems Interconnection Model

| <i>Data Unit</i> | <i>Layer</i> | <i>Function</i> |
|------------------|-----------------|--|
| Data | 7. Application | Network process to application |
| Data | 6. Presentation | Data representation, encryption/decryption, convert dependent data to machine independent data |
| Data | 5. Session | Interhost communication, managing sessions between applications |
| Segments | 4. Transport | End-to-end connections, reliability, and flow control |
| Packet Datagram | 3. Network | Path determination and logical addressing |

OSI Layers

Open Systems Interconnection Model

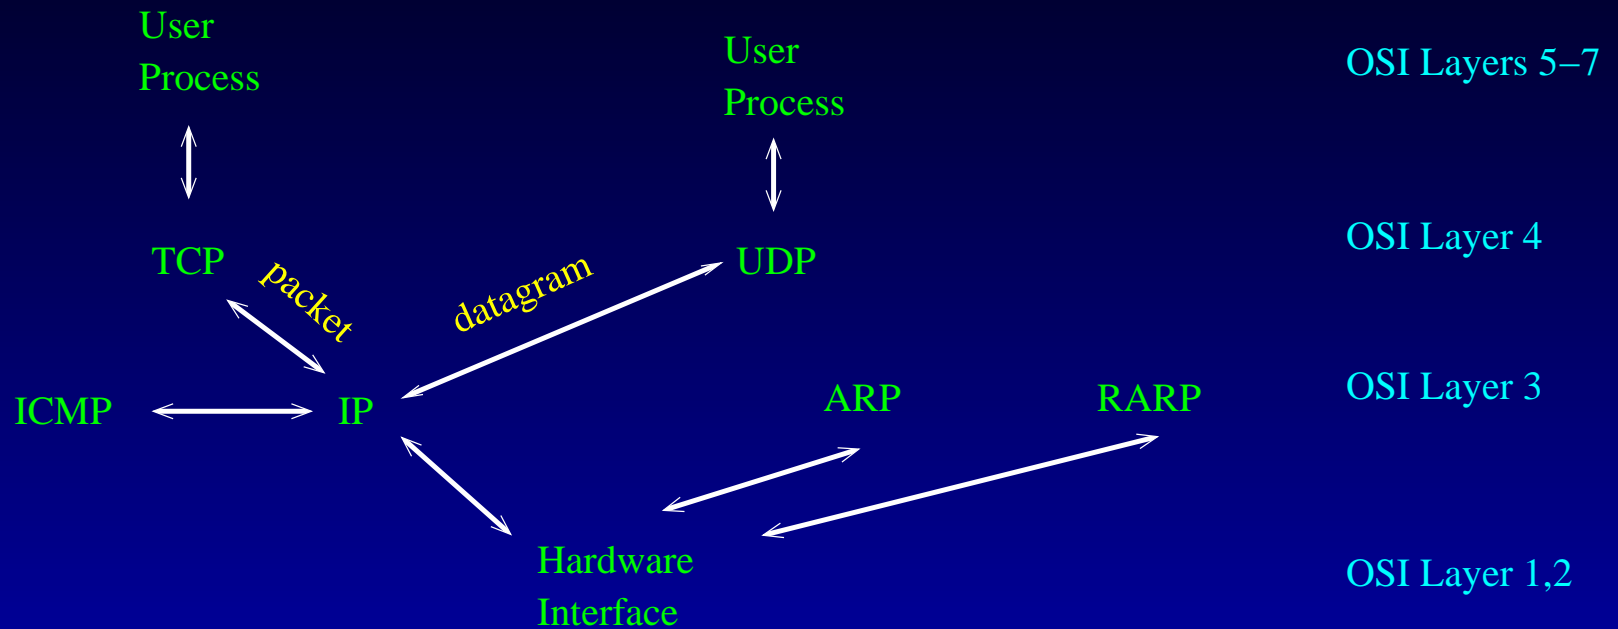
| <i>Data Unit</i> | <i>Layer</i> | <i>Function</i> |
|------------------|-----------------|--|
| Data | 7. Application | Network process to application |
| Data | 6. Presentation | Data representation, encryption/decryption, convert dependent data to machine independent data |
| Data | 5. Session | Interhost communication, managing sessions between applications |
| Segments | 4. Transport | End-to-end connections, reliability, and flow control |
| Packet Datagram | 3. Network | Path determination and logical addressing |
| Frame | 2. Data link | Physical addressing |

OSI Layers

Open Systems Interconnection Model

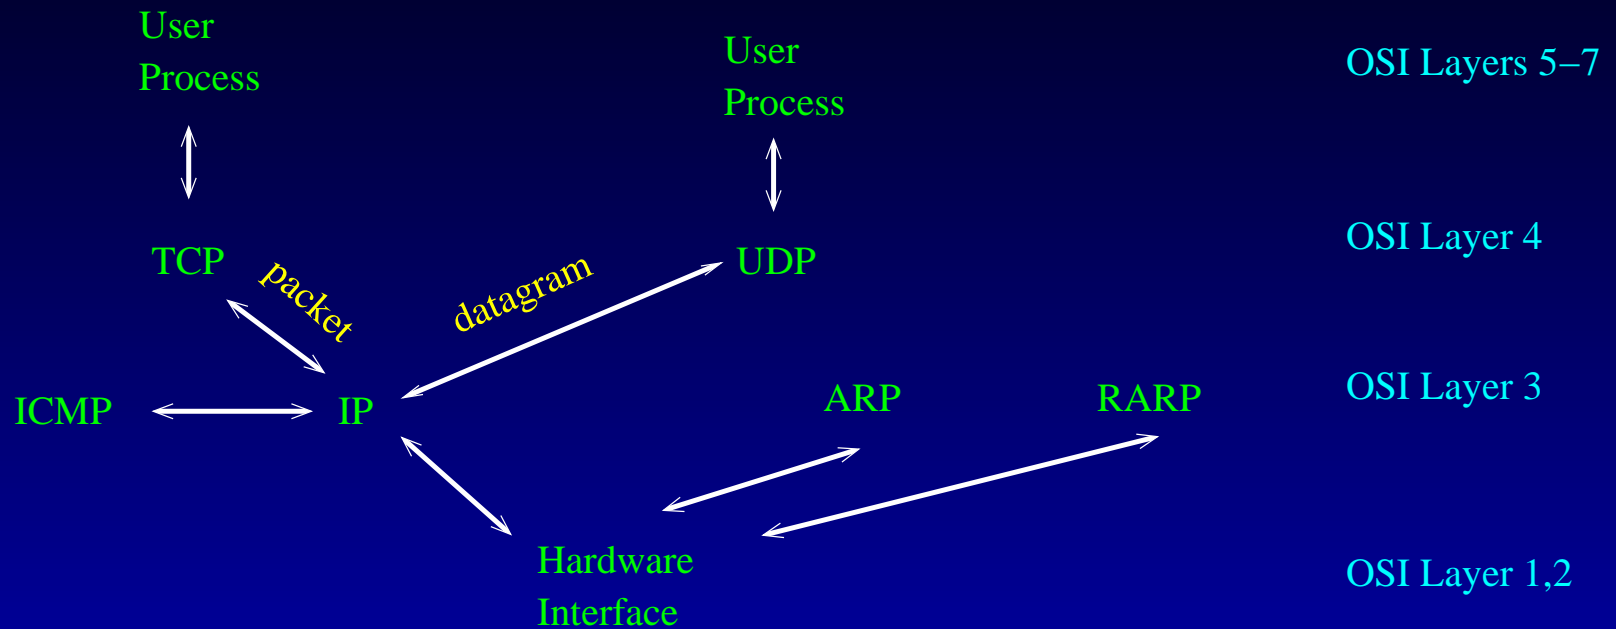
| <i>Data Unit</i> | <i>Layer</i> | <i>Function</i> |
|------------------|-----------------|--|
| Data | 7. Application | Network process to application |
| Data | 6. Presentation | Data representation, encryption/decryption, convert dependent data to machine independent data |
| Data | 5. Session | Interhost communication, managing sessions between applications |
| Segments | 4. Transport | End-to-end connections, reliability, and flow control |
| Packet Datagram | 3. Network | Path determination and logical addressing |
| Frame | 2. Data link | Physical addressing |
| Bit | 1. Physical | Media, signal, and binary transmission |

Layers and Protocols



- IP handles routing. May fragment a packet into multiple packets if the next network can't handle the packet's size, but it will be re-assembled upon delivery

Layers and Protocols



- IP handles routing. May fragment a packet into multiple packets if the next network can't handle the packet's size, but it will be re-assembled upon delivery
- If IP packets arrive so fast that they are discarded, the IP module will send a quench ICMP to the source

Transport Layer

| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|----------------|-----------|------------|------------|
|----------------|-----------|------------|------------|

Transport Layer

| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|---------------------|-----------|------------|------------|
| connection-oriented | | | ✓ |

Transport Layer

| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|---------------------|-----------|------------|------------|
| connection-oriented | | | ✓ |
| message boundaries | ✓ | ✓ | |

Transport Layer

| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|---------------------|-----------|------------|------------|
| connection-oriented | | | ✓ |
| message boundaries | ✓ | ✓ | |
| data checksum | | optional | ✓ |

Transport Layer

| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|---------------------|-----------|------------|------------|
| connection-oriented | | | ✓ |
| message boundaries | ✓ | ✓ | |
| data checksum | | optional | ✓ |
| + acknowledge | | | ✓ |

Transport Layer

| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|------------------------|-----------|------------|------------|
| connection-oriented | | | ✓ |
| message boundaries | ✓ | ✓ | |
| data checksum | | optional | ✓ |
| + acknowledge | | | ✓ |
| timeout to re-transmit | | | ✓ |

Transport Layer

| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|------------------------|-----------|------------|------------|
| connection-oriented | | | ✓ |
| message boundaries | ✓ | ✓ | |
| data checksum | | optional | ✓ |
| + acknowledge | | | ✓ |
| timeout to re-transmit | | | ✓ |
| duplicate detection | | | ✓ |

Transport Layer

| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|------------------------|-----------|------------|------------|
| connection-oriented | | | ✓ |
| message boundaries | ✓ | ✓ | |
| data checksum | | optional | ✓ |
| + acknowledge | | | ✓ |
| timeout to re-transmit | | | ✓ |
| duplicate detection | | | ✓ |
| sequencing | | | ✓ |

Transport Layer

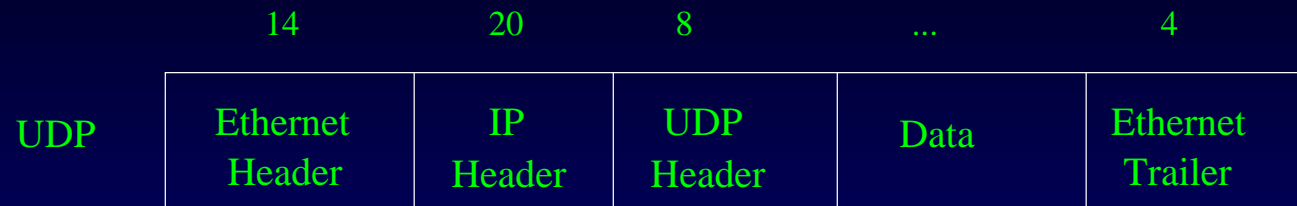
| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|------------------------|-----------|------------|------------|
| connection-oriented | | | ✓ |
| message boundaries | ✓ | ✓ | |
| data checksum | | optional | ✓ |
| + acknowledge | | | ✓ |
| timeout to re-transmit | | | ✓ |
| duplicate detection | | | ✓ |
| sequencing | | | ✓ |
| flow control | | | ✓ |

Transport Layer

| <i>Feature</i> | <i>IP</i> | <i>UDP</i> | <i>TCP</i> |
|------------------------|-----------|------------|------------|
| connection-oriented | | | ✓ |
| message boundaries | ✓ | ✓ | |
| data checksum | | optional | ✓ |
| + acknowledge | | | ✓ |
| timeout to re-transmit | | | ✓ |
| duplicate detection | | | ✓ |
| sequencing | | | ✓ |
| flow control | | | ✓ |

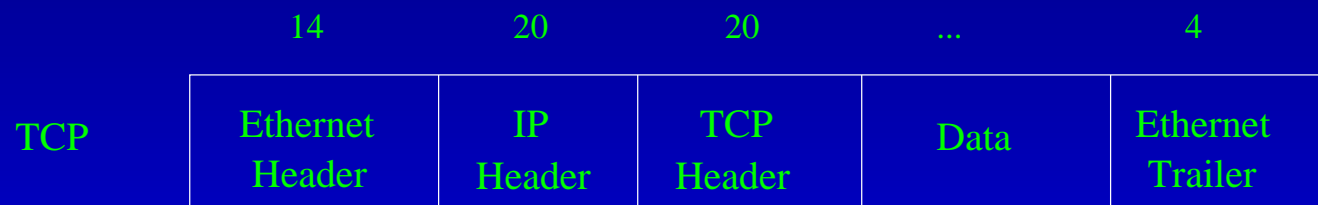
UDP provides port numbers and an optional checksum over IP

Data Packaging



Frame type=IP
 Ethernet 48-bit source address
 Ethernet 48-bit destination address
 Protocol=UDP
 Inet source address
 Inet destination address
 Source Port
 Destination Port

user
 udp
 ip



Protocol=TCP
 Inet 32 bit source address
 Inet 32 bit destination address

user
 tcp
 ip

Clearly, transmission of user data is most efficient when the data section is as large as possible.

Sockets: bind()

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

(1-tuple → 3-tuple)

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- When a socket is created with `socket()`, it has an address family (name space) but no address assigned to it.

Sockets: bind()

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

(1-tuple → 3-tuple)

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- When a socket is created with `socket()`, it has an address family (name space) but no address assigned to it.
- Bind assigns the address specified by `addr`

Sockets: bind()

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

(1-tuple → 3-tuple)

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- When a socket is created with `socket()`, it has an address family (name space) but no address assigned to it.
- Bind assigns the address specified by `addr`
- You'll need to use `accept()` to assign a local address for a full 5-tuple *(for servers)*

Sockets: bind()

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

(1-tuple → 3-tuple)

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- When a socket is created with `socket()`, it has an address family (name space) but no address assigned to it.
- Bind assigns the address specified by `addr`
- You'll need to use `accept()` to assign a local address for a full 5-tuple *(for servers)*
- For `addrlen`: use `sizeof()`

Sockets: bind()

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

(1-tuple → 3-tuple)

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- When a socket is created with `socket()`, it has an address family (name space) but no address assigned to it.
- Bind assigns the address specified by `addr`
- You'll need to use `accept()` to assign a local address for a full 5-tuple *(for servers)*
- For `addrlen`: use `sizeof()`
- server: `sin_addr.s_addr= htonl(INADDR_ANY);`

Sockets: bind()

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

(1-tuple → 3-tuple)

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- When a socket is created with `socket()`, it has an address family (name space) but no address assigned to it.
- Bind assigns the address specified by `addr`
- You'll need to use `accept()` to assign a local address for a full 5-tuple *(for servers)*
- For `addrlen`: use `sizeof()`
- server: `sin_addr.s_addr= htonl(INADDR_ANY);`
- client: `sin_addr= (((struct in_addr *) gethostbyname(...))->h_addr);` *(obsolete - see getaddrinfo())*

Sockets: `bind()`, con't.

- For Unix-style (*local*) addr:

```
struct sockaddr {  
    short  sun_family; // AF_UNIX  
    char   sun_path[]; // unix pathname  
}
```

- For Internet-style addr:

```
struct sockaddr_in {  
    short      sin_family; // AF_INET  
    u_short    port;       // port number  
    struct in_addr sin_addr;  
}
```

Use a cast to pass `in_addr` structure pointers to `bind()`:

```
(struct sockaddr *)&sockaddrin
```

Sockets: listen()

int listen(int sockfd, int backlog);

- sockfd: socket descriptor as returned by socket()

Sockets: listen()

int listen(int sockfd, int backlog);

- sockfd: socket descriptor as returned by socket()
- backlog: qty of clients allowed to be awaiting connection.
Usually maximum is 5. Often ignored.

Sockets: listen()

int listen(int sockfd, int backlog);

- sockfd: socket descriptor as returned by socket()
- backlog: qty of clients allowed to be awaiting connection.
Usually maximum is 5. Often ignored.
- Establishes a connection-oriented socket (SOCK_STREAM or SEQPACKET)

Sockets: listen()

int listen(int sockfd, int backlog);

- sockfd: socket descriptor as returned by socket()
- backlog: qty of clients allowed to be awaiting connection.
Usually maximum is 5. Often ignored.
- Establishes a connection-oriented socket (**SOCK_STREAM** or **SEQPACKET**)
- The socket referred to by sockfd is a passive socket; ie. it needs accept() to complete the 5-tuple

Sockets: listen()

int listen(int sockfd, int backlog);

- sockfd: socket descriptor as returned by socket()
- backlog: qty of clients allowed to be awaiting connection.
Usually maximum is 5. Often ignored.
- Establishes a connection-oriented socket (**SOCK_STREAM** or **SEQPACKET**)
- The socket referred to by sockfd is a passive socket; ie. it needs accept() to complete the 5-tuple
- If a connection request arrives, but the backlog queue is full, the client may receive a -1/ECONNREFUSED error.

sockaddr

```
struct sockaddr {
    unsigned short sa_family;    // address family , AF_XXX
    char           sa_data[14];  // 14 bytes of protocol address
}

struct in_addr {                // IPv4 sockets
    unsigned long  s_addr;      // load this with inet_pton()
}

struct sockaddr_in6 {           // IPv6 sockets
    u_int16_t      sin6_family; // address family , AF_INET6
    u_int16_t      sin6_port;    // port number, in network byte order
    u_int32_t      sin6_scope_id; // scope id
}

struct in6_addr {              // IPv6 address
    unsigned char  s6_addr[16]; // load with inet_pton()
}
```

Many socket functions use one of the above structures cast to a pointer to struct sockaddr.

Utilities: inet_pton()

```
#include <arpa/inet.h>
```

```
int inet_pton(int af, const char *src, void *dst);
```

This function may be used to set up components in various sockaddr family structures.

Utilities: inet_pton()

```
#include <arpa/inet.h>
```

```
int inet_pton(int af, const char *src, void *dst);
```

This function may be used to set up components in various sockaddr family structures. It converts `src` into a network address structure, and then copies that structure to `dst`.

af Address family specification:

Utilities: inet_pton()

```
#include <arpa/inet.h>
```

```
int inet_pton(int af, const char *src, void *dst);
```

This function may be used to set up components in various sockaddr family structures. It converts `src` into a network address structure, and then copies that structure to `dst`.

af Address family specification:

AF_INET Converts an IPv4 dotted-decimal format, “ddd.ddd.ddd.ddd” into a struct `in_addr`.

Utilities: inet_pton()

```
#include <arpa/inet.h>
```

```
int inet_pton(int af, const char *src, void *dst);
```

This function may be used to set up components in various sockaddr family structures. It converts `src` into a network address structure, and then copies that structure to `dst`.

af Address family specification:

AF_INET Converts an IPv4 dotted-decimal format, “ddd.ddd.ddd.ddd” into a struct `in_addr`.

AF_INET6 Converts an IPv6 colon-hexadecimal format, “x:x:x:x:x:x:x:x” into struct `in6_addr`.

Utilities: inet_pton()

#include <arpa/inet.h>

int inet_pton(int af, const char *src, void *dst);

This function may be used to set up components in various sockaddr family structures. It converts `src` into a network address structure, and then copies that structure to `dst`.

af Address family specification:

AF_INET Converts an IPv4 dotted-decimal format, “ddd.ddd.ddd.ddd” into a struct `in_addr`.

AF_INET6 Converts an IPv6 colon-hexadecimal format, “x:x:x:x:x:x:x:x” into struct `in6_addr`.

src A source string describing a numerically-based network address

Utilities: inet_pton()

```
#include <arpa/inet.h>
```

```
int inet_pton(int af, const char *src, void *dst);
```

This function may be used to set up components in various sockaddr family structures. It converts `src` into a network address structure, and then copies that structure to `dst`.

af Address family specification:

AF_INET Converts an IPv4 dotted-decimal format, “ddd.ddd.ddd.ddd” into a `struct in_addr`.

AF_INET6 Converts an IPv6 colon-hexadecimal format, “x:x:x:x:x:x:x:x” into `struct in6_addr`.

src A source string describing a numerically-based network address

dst Destination structure address; for `AF_INET`: must be `sizeof(struct in_addr)` bytes long (4 bytes).

For `AF_INET6`: must be `sizeof(struct in6_addr)` bytes long (16 bytes)

Sockets: connect()

(1-tuple → 5-tuple)

```
int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- sockfd: socket descriptor

Sockets: connect()

(1-tuple → 5-tuple)

int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);

- sockfd: socket descriptor
- addr: Use either a struct sockaddr * or a struct sockaddr_in *.

Sockets: connect()

(1-tuple → 5-tuple)

int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);

- sockfd: socket descriptor
- addr: Use either a struct sockaddr * or a struct sockaddr_in *.
- addrlen: use sizeof(*addr)

Sockets: connect()

(1-tuple → 5-tuple)

int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);

- sockfd: socket descriptor
- addr: Use either a struct sockaddr * or a struct sockaddr_in *.
- addrlen: use sizeof(*addr)
- Connection-oriented clients connect via this call, once. May block.

Sockets: connect()

(1-tuple → 5-tuple)

int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);

- sockfd: socket descriptor
- addr: Use either a struct sockaddr * or a struct sockaddr_in *.
- addrlen: use sizeof(*addr)
- Connection-oriented clients connect via this call, once. May block.
- If the socket sockfd is of type SOCK_DGRAM, then the addr becomes the default send-to address.

Sockets: connect()

(1-tuple → 5-tuple)

int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);

- sockfd: socket descriptor
- addr: Use either a struct sockaddr * or a struct sockaddr_in *.
- addrlen: use sizeof(*addr)
- Connection-oriented clients connect via this call, once. May block.
- If the socket sockfd is of type SOCK_DGRAM, then the addr becomes the default send-to address.
- Connectionless sockets may use this call multiple times.

Sockets: connect()

(1-tuple → 5-tuple)

int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);

- sockfd: socket descriptor
- addr: Use either a struct sockaddr * or a struct sockaddr_in *.
- addrlen: use sizeof(*addr)
- Connection-oriented clients connect via this call, once. May block.
- If the socket sockfd is of type SOCK_DGRAM, then the addr becomes the default send-to address.
- Connectionless sockets may use this call multiple times.
- Connectionless sockets may use sa_family set to AF_UNSPEC to dissolve the association.

Sockets: `accept()`

(3-tuple → 5-tuple)

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

- Used with connection-based sockets (servers)

Sockets: `accept()`

(3-tuple → 5-tuple)

`int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);`

- Used with connection-based sockets (servers)
- Extracts first connection request waiting on the listening queue

Sockets: `accept()`

(3-tuple → 5-tuple)

`int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);`

- Used with connection-based sockets (servers)
- Extracts first connection request waiting on the listening queue
- The call creates a new connected socket, returning a file descriptor for that socket

Sockets: `accept()`

(3-tuple \rightarrow 5-tuple)

`int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);`

- Used with connection-based sockets (servers)
- Extracts first connection request waiting on the listening queue
- The call creates a new connected socket, returning a file descriptor for that socket
- The input `sockfd` is unaffected by the call, and may be used repeatedly

Sockets: `accept()`

(3-tuple → 5-tuple)

`int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);`

- Used with connection-based sockets (servers)
- Extracts first connection request waiting on the listening queue
- The call creates a new connected socket, returning a file descriptor for that socket
- The input `sockfd` is unaffected by the call, and may be used repeatedly
- `addr`: use `struct sockaddr *` or `struct sockaddr_in *`, as appropriate (cast the latter)

Sockets: `accept()`

(3-tuple → 5-tuple)

`int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);`

- Used with connection-based sockets (servers)
- Extracts first connection request waiting on the listening queue
- The call creates a new connected socket, returning a file descriptor for that socket
- The input `sockfd` is unaffected by the call, and may be used repeatedly
- `addr`: use `struct sockaddr *` or `struct sockaddr_in *`, as appropriate (cast the latter)
- Presumes that `sockaddr` had been created with `socket()`, bound to a local address with `bind()`, and `listen()` has been used.

Sockets: `accept()`

(3-tuple \rightarrow 5-tuple)

`int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);`

- Used with connection-based sockets (servers)
- Extracts first connection request waiting on the listening queue
- The call creates a new connected socket, returning a file descriptor for that socket
- The input `sockfd` is unaffected by the call, and may be used repeatedly
- `addr`: use `struct sockaddr *` or `struct sockaddr_in *`, as appropriate (cast the latter)
- Presumes that `sockaddr` had been created with `socket()`, bound to a local address with `bind()`, and `listen()` has been used.
- If no pending connections are on the queue, and `sockaddr` is *not* non-blocking, `accept()` will block.

Sockets: `accept()`

(3-tuple → 5-tuple)

`int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);`

- Used with connection-based sockets (servers)
- Extracts first connection request waiting on the listening queue
- The call creates a new connected socket, returning a file descriptor for that socket
- The input `sockfd` is unaffected by the call, and may be used repeatedly
- `addr`: use `struct sockaddr *` or `struct sockaddr_in *`, as appropriate (cast the latter)
- Presumes that `sockaddr` had been created with `socket()`, bound to a local address with `bind()`, and `listen()` has been used.
- If no pending connections are on the queue, and `sockaddr` is *not* non-blocking, `accept()` will block.
- If no pending connections are on the queue, but `sockaddr` is non-blocking, `accept()` will return a `-1/EAGAIN`.

Sockets: send()

```
ssize_t send(int sockfd, const void *buf, size_t len, int flags);
```

sockfd socket descriptor

Sockets: send()

```
ssize_t send(int sockfd, const void *buf, size_t len, int flags);
```

sockfd socket descriptor

buf buffer of bytes to be sent

Sockets: send()

```
ssize_t send(int sockfd, const void *buf, size_t len, int flags);
```

sockfd socket descriptor

buf buffer of bytes to be sent

len qty of bytes in buffer to be sent

Sockets: send()

ssize_t send(int sockfd, const void *buf, size_t len, int flags);

sockfd socket descriptor

buf buffer of bytes to be sent

len qty of bytes in buffer to be sent

flags 0=normal; MSG_OOB: out-of-band (high priority)

Sockets: send()

ssize_t send(int sockfd, const void *buf, size_t len, int flags);

sockfd socket descriptor

buf buffer of bytes to be sent

len qty of bytes in buffer to be sent

flags 0=normal; MSG_OOB: out-of-band (high priority)

- The send() call expects the sockfd to be in an already connected state (ie. the recipient is known).

Sockets: send()

ssize_t send(int sockfd, const void *buf, size_t len, int flags);

sockfd socket descriptor

buf buffer of bytes to be sent

len qty of bytes in buffer to be sent

flags 0=normal; MSG_OOB: out-of-band (high priority)

- The send() call expects the sockfd to be in an already connected state (ie. the recipient is known).
- Returns -1 on failure; otherwise, a positive number is the qty of bytes actually sent. *may be less than len!*

Sockets: send()

ssize_t send(int sockfd, const void *buf, size_t len, int flags);

sockfd socket descriptor

buf buffer of bytes to be sent

len qty of bytes in buffer to be sent

flags 0=normal; MSG_OOB: out-of-band (high priority)

- The send() call expects the sockfd to be in an already connected state (ie. the recipient is known).
- Returns -1 on failure; otherwise, a positive number is the qty of bytes actually sent. *may be less than len!*
- Equivalent to sendto(sockfd,buf,len,flags,NULL,0) (see next slide)

Sockets: sendto()

```
ssize_t sendto(int sockfd, const void *buf, size_t len, int flags, const struct  
sockaddr *dest_addr, socklen_t addrlen);
```

sockfd...flags same as send()'s

Sockets: sendto()

```
ssize_t sendto(int sockfd, const void *buf, size_t len, int flags, const struct  
sockaddr *dest_addr, socklen_t addrlen);
```

sockfd...flags same as send()'s

dest_addr specifies the destination

Sockets: sendto()

```
ssize_t sendto(int sockfd, const void *buf, size_t len, int flags, const struct  
sockaddr *dest_addr, socklen_t addrlen);
```

sockfd...flags same as send()'s

dest_addr specifies the destination

addrlen specifies the size of dest_addr

Sockets: sendto()

```
ssize_t sendto(int sockfd, const void *buf, size_t len, int flags, const struct  
sockaddr *dest_addr, socklen_t addrlen);
```

sockfd...flags same as send()'s

dest_addr specifies the destination

addrlen specifies the size of dest_addr

- If sendto() is used on a connection-based socket (SOCK_STREAM, SOCK_SEQPACKET), then the last two arguments are ignored.

Sockets: sendto()

ssize_t sendto(int sockfd, const void *buf, size_t len, int flags, const struct sockaddr *dest_addr, socklen_t addrlen);

sockfd...flags same as send()'s

dest_addr specifies the destination

addrlen specifies the size of dest_addr

- If sendto() is used on a connection-based socket (SOCK_STREAM, SOCK_SEQPACKET), then the last two arguments are ignored.
- Otherwise, for connection-less sockets such as SOCK_DGRAM (**datagram**), the recipient's address is specified by dest_addr and addrlen.

Sockets: sendto()

ssize_t sendto(int sockfd, const void *buf, size_t len, int flags, const struct sockaddr *dest_addr, socklen_t addrlen);

sockfd...flags same as send()'s

dest_addr specifies the destination

addrlen specifies the size of dest_addr

- If sendto() is used on a connection-based socket (SOCK_STREAM, SOCK_SEQPACKET), then the last two arguments are ignored.
- Otherwise, for connection-less sockets such as SOCK_DGRAM (datagram), the recipient's address is specified by dest_addr and addrlen.
- If the message is too long, sendto() will return an error EMSGSIZE; the message is not transmitted.

Sockets: `recv()`

`ssize_t recv(int sockfd, void *buf, size_t len, int flags);`

`sockfd` the socket descriptor for a connected socket

Sockets: `recv()`

`ssize_t recv(int sockfd, void *buf, size_t len, int flags);`

`sockfd` the socket descriptor for a connected socket

`buf, len` receives data in this buffer, up to `len` bytes

Sockets: `recv()`

`ssize_t recv(int sockfd, void *buf, size_t len, int flags);`

`sockfd` the socket descriptor for a connected socket

`buf, len` receives data in this buffer, up to `len` bytes

`flags`

Sockets: `recv()`

`ssize_t recv(int sockfd, void *buf, size_t len, int flags);`

`sockfd` the socket descriptor for a connected socket

`buf, len` receives data in this buffer, up to `len` bytes

`flags`

`0` normal

Sockets: `recv()`

`ssize_t recv(int sockfd, void *buf, size_t len, int flags);`

`sockfd` the socket descriptor for a connected socket

`buf, len` receives data in this buffer, up to `len` bytes

`flags`

`0` normal

`MSG_PEEK` allows a peek at the bytes, leaving them for subsequent retrieval
(doesn't remove them from the stream)

Sockets: `recv()`

`ssize_t recv(int sockfd, void *buf, size_t len, int flags);`

`sockfd` the socket descriptor for a connected socket

`buf, len` receives data in this buffer, up to `len` bytes

`flags`

`0` normal

`MSG_PEEK` allows a peek at the bytes, leaving them for subsequent retrieval

(doesn't remove them from the stream)

`MSG_OOB` requests receipt of out-of-band data (which is not normally received)

Sockets: `recv()`

`ssize_t recv(int sockfd, void *buf, size_t len, int flags);`

`sockfd` the socket descriptor for a connected socket

`buf, len` receives data in this buffer, up to `len` bytes

`flags`

`0` normal

`MSG_PEEK` allows a peek at the bytes, leaving them for subsequent retrieval
(doesn't remove them from the stream)

`MSG_OOB` requests receipt of out-of-band data (which is not normally received)

- Will block (unless the socket has been set to non-blocking, see `fcntl()`)

Sockets: `recv()`

`ssize_t recv(int sockfd, void *buf, size_t len, int flags);`

`sockfd` the socket descriptor for a connected socket

`buf, len` receives data in this buffer, up to `len` bytes

`flags`

`0` normal

`MSG_PEEK` allows a peek at the bytes, leaving them for subsequent retrieval
(doesn't remove them from the stream)

`MSG_OOB` requests receipt of out-of-band data (which is not normally received)

- Will block (unless the socket has been set to non-blocking, see `fcntl()`)
- If nonblocking, a -1 will be returned with `errno` set to `EAGAIN`.

Sockets: `recv()`

`ssize_t recv(int sockfd, void *buf, size_t len, int flags);`

`sockfd` the socket descriptor for a connected socket

`buf, len` receives data in this buffer, up to `len` bytes

`flags`

`0` normal

`MSG_PEEK` allows a peek at the bytes, leaving them for subsequent retrieval
(doesn't remove them from the stream)

`MSG_OOB` requests receipt of out-of-band data (which is not normally received)

- Will block (unless the socket has been set to non-blocking, see `fcntl()`)
- If nonblocking, a -1 will be returned with `errno` set to `EAGAIN`.
- Returns -1 on failure, or a positive qty of bytes actually received (which may be less than those sent!)

Sockets: recvfrom()

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags, struct sockaddr  
*src_addr, socklen_t *addrlen);
```

sid...flag just like recv()'s

Sockets: recvfrom()

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags, struct sockaddr  
*src_addr, socklen_t *addrlen);
```

src_addr, addrlen just like `recv()`'s

src_addr, addrlen specifies the sender's address.

Sockets: recvfrom()

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags, struct sockaddr  
*src_addr, socklen_t *addrlen);
```

src_addr, addrlen just like `recv()`'s

src_addr, addrlen specifies the sender's address.

- Returns like `recv()`

Sockets: recvfrom()

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags, struct sockaddr  
*src_addr, socklen_t *addrlen);
```

src_addr, addrlen just like recv()'s

src_addr, addrlen specifies the sender's address.

- Returns like recv()
- Primarily for SOCK_DGRAM (datagrams)

Sockets: shutdown()

```
int shutdown(int sockfd, int how);
```

sockfd socket descriptor to be shut down

Sockets: shutdown()

int shutdown(int sockfd, int how);

sockfd socket descriptor to be shut down

how controls partial shutdowns

Sockets: shutdown()

int shutdown(int sockfd, int how);

sockfd socket descriptor to be shut down

how controls partial shutdowns

SHUT_RD further reading (receptions) disallowed

Sockets: shutdown()

int shutdown(int sockfd, int how);

sockfd socket descriptor to be shut down

how controls partial shutdowns

SHUT_RD further reading (receptions) disallowed

SHUT_WR further writing (transmissions) disallowed

Sockets: shutdown()

int shutdown(int sockfd, int how);

sockfd socket descriptor to be shut down

how controls partial shutdowns

SHUT_RD further reading (receptions) disallowed

SHUT_WR further writing (transmissions) disallowed

SHUT_RDWR subsequent reading and writing via sockfd is disallowed

Sockets: close()

```
#include <unistd.h>
```

```
int close(int fd);
```

- This call closes file descriptors, including socket descriptors.

Sockets: close()

```
#include <unistd.h>
```

```
int close(int fd);
```

- This call closes file descriptors, including socket descriptors.
- The descriptor will no longer refer to any file or socket and may be re-used.

Sockets: close()

```
#include <unistd.h>
```

```
int close(int fd);
```

- This call closes file descriptors, including socket descriptors.
- The descriptor will no longer refer to any file or socket and may be re-used.
- If `fd` is the last file descriptor holding open the file, then resources associated with the open file descriptor are freed.

(see connection.c packet.c)

Utilities: gethostbyname()

```
#include <netdb.h>
```

```
extern int h_errno;
```

```
struct hostent *gethostbyname(const char *name);
```

```
#define h_addr h_addr_list[0]
```

name the name of the host, using IPv4 standard dot notation or in IPv6 notation; in either colon or dot notation accepted.

- Returns a struct hostent (*see next slide*)
- A host may have more than one address if it is “multihomed”
(ie. the machine has more than one internet interface, each with its own address)
- Considered to be obsolete; use getaddrinfo() instead

Utilities: The `hostent` Structure

```
struct hostent {  
    char    *h_name;           // official name of host  
    char **h_aliases;         // alias list  
    int     h_length;         // length of address  
    char **h_addr_host;       // list of addresses from name server  
                                // (a NULL terminated list)  
}
```

Socket Code Fragment

```
struct sockaddr_in  sin; // internet socket addressing
struct hostent      *phostent= NULL;

// open a socket for clients
sktfd= socket(AF_INET,SOCK_STREAM,0);

// initialize the sockaddr_in
memset(&sin,0x00,sizeof(struct sockaddr_in));
sin.sin_family = AF_INET;
sin.sin_addr   = *((struct in_addr *) phostent->h_addr;
sin.sin_port   = htons((u_short) port);

// get a hostent structure
phostent       = gethostbyname(hostname);

// connect
if(connect(sktfd,(struct sockaddr*) &sin,sizeof(sin)) < 0){
    perror("");
}
```

Utilities: IP and Name Lookups

| Transformation | Obsolete Functions | New Functions |
|-----------------------|--------------------|---------------|
| name \rightarrow IP | gethostbyname() | getaddrinfo() |
| IP \rightarrow name | gethostbyaddr() | getnameinfo() |

The newer functions can use IPv4 or IPv6 transparently, and they can handle services such as “http” as well as port numbers.

Utilities: getaddrinfo()

```
int getaddrinfo(const char *node, const char *service, const struct addrinfo  
*hints, struct addrinfo **res);
```

```
void freeaddrinfo(struct addrinfo *res);
```

node numerical network address or a hostname string

Utilities: getaddrinfo()

```
int getaddrinfo(const char *node, const char *service, const struct addrinfo  
*hints, struct addrinfo **res);
```

```
void freeaddrinfo(struct addrinfo *res);
```

node numerical network address or a hostname string

service sets the port; service name (see [man 5 services](#)). May be null.

Utilities: getaddrinfo()

```
int getaddrinfo(const char *node, const char *service, const struct addrinfo  
*hints, struct addrinfo **res);
```

```
void freeaddrinfo(struct addrinfo *res);
```

node numerical network address or a hostname string

service sets the port; service name (see [man 5 services](#)). May be null.

hints See struct addrinfo below. Specifies criteria for selecting socket address structure in the list pointed to by res

Utilities: getaddrinfo()

```
int getaddrinfo(const char *node, const char *service, const struct addrinfo  
*hints, struct addrinfo **res);
```

```
void freeaddrinfo(struct addrinfo *res);
```

node numerical network address or a hostname string

service sets the port; service name (see [man 5 services](#)). May be null.

hints See struct addrinfo below. Specifies criteria for selecting socket address structure in the list pointed to by res

res A pointer to the head of a linked list of addrinfo structures. ai_next gives the next link.

Utilities: getaddrinfo()

```
int getaddrinfo(const char *node, const char *service, const struct addrinfo  
*hints, struct addrinfo **res);
```

```
void freeaddrinfo(struct addrinfo *res);
```

node numerical network address or a hostname string

service sets the port; service name (see [man 5 services](#)). May be null.

hints See struct addrinfo below. Specifies criteria for selecting socket address structure in the list pointed to by res

res A pointer to the head of a linked list of addrinfo structures. ai_next gives the next link.

```
struct addrinfo {  
    int             ai_flags;  
    int             ai_family;  
    int             ai_socktype;  
    int             ai_protocol;  
    size_t          ai_addrlen;  
    struct sockaddr * ai_addr;  
    char            * ai_canonname;  
    struct addrinfo * ai_next;  
}
```

see next slide

Utilities: `getaddrinfo()`, con't.

- Identifies an internet host or service

Utilities: `getaddrinfo()`, con't.

- Identifies an internet host or service
- Returns one or more `addrinfo` structures, each of which contains an internet address that can be used in `bind()` or `connect()`.

Utilities: `getaddrinfo()`, con't.

- Identifies an internet host or service
- Returns one or more `addrinfo` structures, each of which contains an internet address that can be used in `bind()` or `connect()`.
- Reentrant

Utilities: `getaddrinfo()`, con't.

- Identifies an internet host or service
- Returns one or more `addrinfo` structures, each of which contains an internet address that can be used in `bind()` or `connect()`.
- Reentrant
- `ai_family`, `ai_socktype`, and `si_protocol` mean the same as with `socket()`

Utilities: `getaddrinfo()`, con't.

- Identifies an internet host or service
- Returns one or more `addrinfo` structures, each of which contains an internet address that can be used in `bind()` or `connect()`.
- Reentrant
- `ai_family`, `ai_socktype`, and `si_protocol` mean the same as with `socket()`
- Pointer to the socket address is placed in `ai_addr`

Utilities: getaddrinfo(), con't.

- Identifies an internet host or service
- Returns one or more addrinfo structures, each of which contains an internet address that can be used in bind() or connect().
- Reentrant
- ai_family, ai_socktype, and si_protocol mean the same as with socket()
- Pointer to the socket address is placed in ai_addr
- Length of socket address (in bytes) is placed in ai_addrlen

Utilities: getaddrinfo(), con't.

- Identifies an internet host or service
- Returns one or more addrinfo structures, each of which contains an internet address that can be used in bind() or connect().
- Reentrant
- ai_family, ai_socktype, and si_protocol mean the same as with socket()
- Pointer to the socket address is placed in ai_addr
- Length of socket address (in bytes) is placed in ai_addrlen
- freeaddrinfo(struct addrinfo *res) frees up the res structure linked list

Utilities: getaddrinfo(), con't.

- Identifies an internet host or service
- Returns one or more addrinfo structures, each of which contains an internet address that can be used in bind() or connect().
- Reentrant
- ai_family, ai_socktype, and si_protocol mean the same as with socket()
- Pointer to the socket address is placed in ai_addr
- Length of socket address (in bytes) is placed in ai_addrlen
- freeaddrinfo(struct addrinfo *res) frees up the res structure linked list

(see using_getaddrinfo.c)

Utilities: getnameinfo()

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t  
hostlen, char *serv, size_t servlen, int flags);
```

sa (input) pointer to a generic socket address structure (either `sockaddr_in` or `sockaddr_in6`)

Utilities: getnameinfo()

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t  
hostlen, char *serv, size_t servlen, int flags);
```

sa (input) pointer to a generic socket address structure (either `sockaddr_in` or `sockaddr_in6`)

salen sa is of size salen.

Utilities: getnameinfo()

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t  
hostlen, char *serv, size_t servlen, int flags);
```

sa (input) pointer to a generic socket address structure (either `sockaddr_in` or `sockaddr_in6`)

salen sa is of size salen.

host Pointer to caller-allocated buffer with a null-terminated host name

Utilities: getnameinfo()

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t  
hostlen, char *serv, size_t servlen, int flags);
```

sa (input) pointer to a generic socket address structure (either `sockaddr_in` or `sockaddr_in6`)

salen sa is of size salen.

host Pointer to caller-allocated buffer with a null-terminated host name

hostlen Length of host buffer

Utilities: getnameinfo()

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t  
hostlen, char *serv, size_t servlen, int flags);
```

sa (input) pointer to a generic socket address structure (either `sockaddr_in` or `sockaddr_in6`)

salen sa is of size salen.

host Pointer to caller-allocated buffer with a null-terminated host name

hostlen Length of host buffer

serv Pointer to caller-allocated buffer with a null-terminated service name

Utilities: getnameinfo()

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t  
hostlen, char *serv, size_t servlen, int flags);
```

sa (input) pointer to a generic socket address structure (either `sockaddr_in` or `sockaddr_in6`)

salen `sa` is of size `salen`.

host Pointer to caller-allocated buffer with a null-terminated host name

hostlen Length of `host` buffer

serv Pointer to caller-allocated buffer with a null-terminated service name

servlen Length of `serv` buffer

Utilities: getnameinfo()

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t  
hostlen, char *serv, size_t servlen, int flags);
```

sa (input) pointer to a generic socket address structure (either `sockaddr_in` or `sockaddr_in6`)

salen `sa` is of size `salen`.

host Pointer to caller-allocated buffer with a null-terminated host name

hostlen Length of `host` buffer

serv Pointer to caller-allocated buffer with a null-terminated service name

servlen Length of `serv` buffer

flags Modifies behavior as follows: (see next slide)

Utilities: getnameinfo()

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t  
hostlen, char *serv, size_t servlen, int flags);
```

sa (input) pointer to a generic socket address structure (either `sockaddr_in` or `sockaddr_in6`)

salen `sa` is of size `salen`.

host Pointer to caller-allocated buffer with a null-terminated host name

hostlen Length of `host` buffer

serv Pointer to caller-allocated buffer with a null-terminated service name

servlen Length of `serv` buffer

flags Modifies behavior as follows: (see next slide)

- `getnameinfo()` converts a socket address into a corresponding host and service, depending on the protocol.

Utilities: getnameinfo()

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t  
hostlen, char *serv, size_t servlen, int flags);
```

sa (input) pointer to a generic socket address structure (either `sockaddr_in` or `sockaddr_in6`)

salen `sa` is of size `salen`.

host Pointer to caller-allocated buffer with a null-terminated host name

hostlen Length of `host` buffer

serv Pointer to caller-allocated buffer with a null-terminated service name

servlen Length of `serv` buffer

flags Modifies behavior as follows: (see next slide)

- `getnameinfo()` converts a socket address into a corresponding host and service, depending on the protocol.
- The function is re-entrant, and allows programs to avoid IPv4 vs IPv6 dependencies

Utilities: getnameinfo(), con't.

flags Modifies behavior of getnameinfo() as follows:

NI_NAMEREQD An error will be returned if the hostname cannot be determined

Utilities: getnameinfo(), con't.

flags Modifies behavior of getnameinfo() as follows:

NI_NAMEREQD An error will be returned if the hostname cannot be determined

NI_DGRAM Service is a datagram (**UDP**). Required for those ports whose associated services support both UDP and TCP.

Utilities: getnameinfo(), con't.

flags Modifies behavior of getnameinfo() as follows:

NI_NAMEREQD An error will be returned if the hostname cannot be determined

NI_DGRAM Service is a datagram (**UDP**). Required for those ports whose associated services support both UDP and TCP.

NI_NOFQDN Return only hostname portion of fully qualified domain name for local hosts

Utilities: getnameinfo(), con't.

flags Modifies behavior of getnameinfo() as follows:

NI_NAMEREQD An error will be returned if the hostname cannot be determined

NI_DGRAM Service is a datagram (**UDP**). Required for those ports whose associated serves support both UDP and TCP.

NI_NOFQDN Return only hostname portion of fully qualified domain name for local hosts

NI_NUMERICHOST returns the numeric form of the hostname

Utilities: getnameinfo(), con't.

flags Modifies behavior of getnameinfo() as follows:

NI_NAMEREQD An error will be returned if the hostname cannot be determined

NI_DGRAM Service is a datagram (**UDP**). Required for those ports whose associated serves support both UDP and TCP.

NI_NOFQDN Return only hostname portion of fully qualified domain name for local hosts

NI_NUMERICHOST returns the numeric form of the hostname

NI_NUMERICSERV returns the numeric form of the service address

Utilities: gethostbyaddr()

struct hostent *gethostbyaddr(const void *addr, socklen_t len, int type);

addr a pointer to a struct depending on the address type (ex. struct in_addr *), which was probably obtained via a call to inet_addr()

Utilities: gethostbyaddr()

struct hostent *gethostbyaddr(const void *addr, socklen_t len, int type);

addr a pointer to a struct depending on the address type (ex. struct in_addr *), which was probably obtained via a call to inet_addr()

len length of addr structure being passed

Utilities: gethostbyaddr()

struct hostent *gethostbyaddr(const void *addr, socklen_t len, int type);

addr a pointer to a struct depending on the address type (ex. struct in_addr *), which was probably obtained via a call to inet_addr()

len length of addr structure being passed

type Valid address types include AF_INET and AF_INET6.

Utilities: gethostbyaddr()

struct hostent *gethostbyaddr(const void *addr, socklen_t len, int type);

addr a pointer to a struct depending on the address type (ex. struct in_addr *), which was probably obtained via a call to inet_addr()

len length of addr structure being passed

type Valid address types include AF_INET and AF_INET6.

- Returns a structure of type hostent (see previous slide)

Utilities: getservbyname()

```
#include <netdb.h>
```

```
struct servent *getservbyname(const char *name, const char *proto);
```

- Returns a servent structure using the service name database (ie. [/etc/services](#))

Utilities: getservbyname()

```
#include <netdb.h>
```

```
struct servent *getservbyname(const char *name, const char *proto);
```

- Returns a servent structure using the service name database (ie. /etc/services)

```
struct servent {  
    char    *    s_name;    official service name  
    char    **   s_aliases;  alias list  
    int      s_port;    port number, already in network byte order  
    char     *    s_proto;    protocol to use  
}
```

Utilities: getservbyname()

#include <netdb.h>

struct servent *getservbyname(const char *name, const char *proto);

- Returns a servent structure using the service name database (ie. /etc/services)

```
struct servent {  
    char    *    s_name;      official service name  
    char    **   s_aliases;   alias list  
    int      s_port;         port number, already in network byte order  
    char     *    s_proto;    protocol to use  
}
```

- To fill in a struct sockaddr_in's short sin_family: use gethostbyname()->h_addrtype

Utilities: getservbyname()

#include <netdb.h>

struct servent *getservbyname(const char *name, const char *proto);

- Returns a servent structure using the service name database (ie. /etc/services)

```
struct servent {  
    char    *    s_name;      official service name  
    char    **   s_aliases;   alias list  
    int      s_port;         port number, already in network byte order  
    char     *    s_proto;    protocol to use  
}
```

- To fill in a struct sockaddr_in's short sin_family: use gethostbyname()->h_addrtype
- To fill in sin_port: gethostbyname(...)->h_addr

Utilities: Byte Ordering

```
#include <arpa/inet.h>
```

```
#include <netinet/in.h> (some require this header file instead)
```

```
uint32_t htonl(uint32_t hostlong);
```

```
uint16_t htons(uint16_t hostshort);
```

```
uint32_t ntohl(uint32_t netlong);
```

```
uint16_t ntohs(uint16_t netshort);
```

- Different machines can have different architectures; in particular, integers can be represented with high-byte first or last.

Utilities: Byte Ordering

```
#include <arpa/inet.h>
```

```
#include <netinet/in.h> (some require this header file instead)
```

```
uint32_t htonl(uint32_t hostlong);
```

```
uint16_t htons(uint16_t hostshort);
```

```
uint32_t ntohl(uint32_t netlong);
```

```
uint16_t ntohs(uint16_t netshort);
```

- Different machines can have different architectures; in particular, integers can be represented with high-byte first or last.
- On Intel based machines, the least significant byte is first; network order has most significant byte first.

Utilities: Byte Ordering

```
#include <arpa/inet.h>
```

```
#include <netinet/in.h> (some require this header file instead)
```

```
uint32_t htonl(uint32_t hostlong);
```

```
uint16_t htons(uint16_t hostshort);
```

```
uint32_t ntohl(uint32_t netlong);
```

```
uint16_t ntohs(uint16_t netshort);
```

- Different machines can have different architectures; in particular, integers can be represented with high-byte first or last.
- On Intel based machines, the least significant byte is first; network order has most significant byte first.
- `htonl()` converts unsigned integers from host format to network byte order

Utilities: Byte Ordering

```
#include <arpa/inet.h>
```

```
#include <netinet/in.h> (some require this header file instead)
```

```
uint32_t htonl(uint32_t hostlong);
```

```
uint16_t htons(uint16_t hostshort);
```

```
uint32_t ntohl(uint32_t netlong);
```

```
uint16_t ntohs(uint16_t netshort);
```

- Different machines can have different architectures; in particular, integers can be represented with high-byte first or last.
- On Intel based machines, the least significant byte is first; network order has most significant byte first.
- `htonl()` converts unsigned integers from host format to network byte order
- `htons()` converts unsigned short integers from host format to network byte order

Utilities: Byte Ordering

```
#include <arpa/inet.h>
```

```
#include <netinet/in.h> (some require this header file instead)
```

```
uint32_t htonl(uint32_t hostlong);
```

```
uint16_t htons(uint16_t hostshort);
```

```
uint32_t ntohl(uint32_t netlong);
```

```
uint16_t ntohs(uint16_t netshort);
```

- Different machines can have different architectures; in particular, integers can be represented with high-byte first or last.
- On Intel based machines, the least significant byte is first; network order has most significant byte first.
- `htonl()` converts unsigned integers from host format to network byte order
- `htons()` converts unsigned short integers from host format to network byte order
- `ntohl()` converts unsigned integers from network byte order to host format

Utilities: Byte Ordering

```
#include <arpa/inet.h>
```

```
#include <netinet/in.h> (some require this header file instead)
```

```
uint32_t htonl(uint32_t hostlong);
```

```
uint16_t htons(uint16_t hostshort);
```

```
uint32_t ntohl(uint32_t netlong);
```

```
uint16_t ntohs(uint16_t netshort);
```

- Different machines can have different architectures; in particular, integers can be represented with high-byte first or last.
- On Intel based machines, the least significant byte is first; network order has most significant byte first.
- `htonl()` converts unsigned integers from host format to network byte order
- `htons()` converts unsigned short integers from host format to network byte order
- `ntohl()` converts unsigned integers from network byte order to host format
- `ntohs()` converts unsigned short integers from network byte order to host format

Utilities: Byte Operations

SysV `#include <string.h>`
`void *memcpy(void *dest, const void *src, size_t n);`
`void *memmove(void *dest, const void *src, size_t n);`
`void *memset(void *s, int c, size_t n);`
`int memcmp(const void *s1, const void *s2, size_t n);`

BSD `#include <strings.h>`
* `void bcopy(const void *src, void *dest, size_t n);`
* `void bzero(void *s, size_t n);`
* `int bcmp(const void *s1, const void *s2, size_t n);`
* *these functions are deprecated*

`memcpy()`, `bcopy()`, `memmove()` copies `n` bytes from `src` to `dest`. Overlapping
dest/target permitted.

Utilities: Byte Operations

SysV **#include <string.h>**
void *memcpy(void *dest, const void *src, size_t n);
void *memmove(void *dest, const void *src, size_t n);
void *memset(void *s, int c, size_t n);
int memcmp(const void *s1, const void *s2, size_t n);

BSD **#include <strings.h>**
* void bcopy(const void *src, void *dest, size_t n);
* **void bzero(void *s, size_t n);**
* int bcmp(const void *s1, const void *s2, size_t n);
* *these functions are deprecated*

memcpy(), bcopy(), memmove() copies n bytes from src to dest. Overlapping dest/target permitted.

memset(), bzero() Sets s to the specified byte (SysV) or to zero (BSD)

Utilities: Byte Operations

SysV **#include <string.h>**
void *memcpy(void *dest, const void *src, size_t n);
void *memmove(void *dest, const void *src, size_t n);
void *memset(void *s, int c, size_t n);
int memcmp(const void *s1, const void *s2, size_t n);

BSD **#include <strings.h>**
* void bcopy(const void *src, void *dest, size_t n);
* void bzero(void *s, size_t n);
* **int bcmp(const void *s1, const void *s2, size_t n);**
* *these functions are deprecated*

memcpy(), bcopy(), memmove() copies *n* bytes from *src* to *dest*. Overlapping *dest/target* permitted.

memset(), bzero() Sets *s* to the specified byte (**SysV**) or to zero (**BSD**)

memcmp(), bcmp() Compares the first *n* bytes of *s1* and *s2*. Returns an integer greater than or less than zero for the first mismatched byte. If *n* bytes are identical, then zero is returned.

Utilities: Address Conversion

```
#include <sys/socket.h>
```

```
#include <netinet/in.h>
```

```
#include <arpa/inet.h>
```

```
int inet_aton(const char *cp, struct in_addr *inp);
```

```
char *inet_ntoa(struct in_addr in);
```

```
in_addr_t inet_network(const char *cp);
```

inet_aton() converts an IPv4 number and dot notation string into a struct in_addr structure.

Returns non-zero if valid, 0 otherwise.

Utilities: Address Conversion

```
#include <sys/socket.h>
```

```
#include <netinet/in.h>
```

```
#include <arpa/inet.h>
```

```
int inet_aton(const char *cp, struct in_addr *inp);
```

```
char *inet_ntoa(struct in_addr in);
```

```
in_addr_t inet_network(const char *cp);
```

inet_aton() converts an IPv4 number and dot notation string into a struct in_addr structure.

Returns non-zero if valid, 0 otherwise.

inet_ntoa() reverses inet_aton()

Utilities: Address Conversion

```
#include <sys/socket.h>
```

```
#include <netinet/in.h>
```

```
#include <arpa/inet.h>
```

```
int inet_aton(const char *cp, struct in_addr *inp);
```

```
char *inet_ntoa(struct in_addr in);
```

```
in_addr_t inet_network(const char *cp);
```

inet_aton() converts an IPv4 number and dot notation string into a struct in_addr structure.

Returns non-zero if valid, 0 otherwise.

inet_ntoa() reverses inet_aton()

inet_network() converts an IPv4 number and dot notation string into a number in network byte order.

Utilities: Address Conversion

```
#include <sys/socket.h>
```

```
#include <netinet/in.h>
```

```
#include <arpa/inet.h>
```

```
int inet_aton(const char *cp, struct in_addr *inp);
```

```
char *inet_ntoa(struct in_addr in);
```

```
in_addr_t inet_network(const char *cp);
```

inet_aton() converts an IPv4 number and dot notation string into a struct in_addr structure.

Returns non-zero if valid, 0 otherwise.

inet_ntoa() reverses inet_aton()

inet_network() converts an IPv4 number and dot notation string into a number in network byte order.

struct in_addr inet_network() returns the following structure:

```
typedef uint32_t in_addr_t;  
struct in_addr {  
    in_addr_t s_addr;  
};
```

Utilities: getpeername()

```
#include <sys/socket.h>
```

```
int getpeername(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

sockfd an open socket descriptor

Utilities: getpeername()

```
#include <sys/socket.h>
```

```
int getpeername(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

sockfd an open socket descriptor

struct sockaddr *addr getpeername() returns the address of the peer connected to the socket sockfd.

Utilities: getpeername()

#include <sys/socket.h>

int getpeername(int sockfd, struct sockaddr *addr, **socklen_t** *addrlen);

sockfd an open socket descriptor

struct sockaddr *addr getpeername() returns the address of the peer connected to the socket sockfd.

addrlen User should initialize this to the actual size of the addr buffer being passed. On return it holds the actual size of the name (in bytes). The name will be truncated if the buffer is too small.

Utilities: getpeername()

#include <sys/socket.h>

int getpeername(int sockfd, struct sockaddr *addr, socklen_t *addrlen);

sockfd an open socket descriptor

struct sockaddr *addr getpeername() returns the address of the peer connected to the socket sockfd.

addrlen User should initialize this to the actual size of the addr buffer being passed. On return it holds the actual size of the name (in bytes). The name will be truncated if the buffer is too small.

The getpeername() function converts a socket file descriptor into a sockaddr-based peername.

Utilities: gethostname()

```
#include <unistd.h>
```

```
int gethostname(char *name, size_t len);
```

- This function returns the null-byte terminated hostname in the name buffer, which has a length of len bytes.

Utilities: gethostname()

```
#include <unistd.h>
```

```
int gethostname(char *name, size_t len);
```

- This function returns the null-byte terminated hostname in the name buffer, which has a length of len bytes.
- The hostname will be truncated if len isn't large enough.

Utilities: gethostname()

```
#include <unistd.h>
```

```
int gethostname(char *name, size_t len);
```

- This function returns the null-byte terminated hostname in the `name` buffer, which has a length of `len` bytes.
- The hostname will be truncated if `len` isn't large enough.
- It may or may not have a terminating null byte in such a case.

Ports

Reserved ports

1 -1023

Ports

| | |
|--------------------------------|-----------|
| Reserved ports | 1 -1023 |
| Automatically assignable ports | 1024-5000 |

Ports

| | |
|--------------------------------|-----------|
| Reserved ports | 1 -1023 |
| Automatically assignable ports | 1024-5000 |
| Ports assigned by rresvport() | 512 -1023 |

Ports

| | |
|--|-----------|
| Reserved ports | 1 -1023 |
| Automatically assignable ports | 1024-5000 |
| Ports assigned by <code>rresvport()</code> | 512 -1023 |

- `rresvport(int *port)` returns a descriptor to a socket with an address in the privileged port space.

Only the superuser privileged processes are allowed to bind an address with this sort of port to a socket.

Ports

| | |
|--------------------------------|-----------|
| Reserved ports | 1 -1023 |
| Automatically assignable ports | 1024-5000 |
| Ports assigned by rresvport() | 512 -1023 |

- `rresvport(int *port)` returns a descriptor to a socket with an address in the privileged port space.
Only the superuser privileged processes are allowed to bind an address with this sort of port to a socket.
- See `/etc/services` for a list of service names, ports, and protocol (ie. whether it uses udp or tcp)