

# Simple Sockets Library

- You may obtain this library from my website: <http://www.drchip.org/astronaut/ssl>

# Simple Sockets Library

- You may obtain this library from my website: <http://www.drchip.org/astronaut/ssl>
- Written to make using connection-oriented (tcp) sockets simple

# Simple Sockets Library

- You may obtain this library from my website: <http://www.drchip.org/astronaut/ssl>
- Written to make using connection-oriented (tcp) sockets simple
- Written to resemble C's file i/o:

*Sclose() Sgets() Sopen() Sprintf() Sputs() Sread() Sscanf() Swrite() Svprintf()*

# Simple Sockets Library

- You may obtain this library from my website: <http://www.drchip.org/astronaut/ssl>
- Written to make using connection-oriented (tcp) sockets simple
- Written to resemble C's file i/o:  
*Sclose() Sgets() Sopen() Sprintf() Sputs() Sread() Sscanf() Swrite() Svprintf()*
- Uses automatically assigned ports, so if a server can be brought up, it will be

# Simple Sockets Library

- You may obtain this library from my website: <http://www.drchip.org/astronaut/ssl>
- Written to make using connection-oriented (tcp) sockets simple
- Written to resemble C's file i/o:  
*Sclose() Sgets() Sopen() Sprintf() Sputs() Sread() Sscanf() Swrite() Svprintf()*
- Uses automatically assigned ports, so if a server can be brought up, it will be
- Servers have names, not fixed ports (solves REUSEADDR problems, linger setting, zombie sockets)

# SSL: Sopen()

**#include "sockets.h"**

**Socket \*Sopen(char \*skthost, char \*mode)**

**skthost** Servers: name of socket.

Clients: name of socket @ hostname

# SSL: Sopen()

**#include "sockets.h"**

**Socket \*Sopen(char \*skthost, char \*mode)**

**skthost** Servers: name of socket.

Clients: name of socket @ hostname

**mode** Tells Sopen the type of socket being opened

# SSL: Sopen()

**#include "sockets.h"**

**Socket \*Sopen(char \*skthost, char \*mode)**

**skthost** Servers: name of socket.

Clients: name of socket @ hostname

**mode** Tells Sopen the type of socket being opened

s open a server



# SSL: Sopen()

**#include “sockets.h”**

**Socket \*Sopen(char \*skthost, char \*mode)**

**skthost** Servers: name of socket.

Clients: name of socket @ hostname

**mode** Tells Sopen the type of socket being opened

**s** open a server

**S** open a server (but, on failure, retry)

# SSL: Sopen()

**#include “sockets.h”**

**Socket \*Sopen(char \*skthost, char \*mode)**

**skthost** Servers: name of socket.

Clients: name of socket @ hostname

**mode** Tells Sopen the type of socket being opened

**s** open a server

**S** open a server (but, on failure, retry)

**c** open a client

# SSL: Sopen()

**#include "sockets.h"**

**Socket \*Sopen(char \*skthost, char \*mode)**

**skthost** Servers: name of socket.

Clients: name of socket @ hostname

**mode** Tells Sopen the type of socket being opened

**s** open a server

**S** open a server (but, on failure, retry)

**c** open a client

**s###** open a server with the specified port number

# SSL: Sopen()

**#include “sockets.h”**

**Socket \*Sopen(char \*skthost, char \*mode)**

**skthost** Servers: name of socket.

Clients: name of socket @ hostname

**mode** Tells Sopen the type of socket being opened

**s** open a server

**S** open a server (but, on failure, retry)

**c** open a client

**s###** open a server with the specified port number

**c###** open a client to a server with the specified port number

# SSL: Sopen()

**#include “sockets.h”**

**Socket \*Sopen(char \*skthost, char \*mode)**

**skthost** Servers: name of socket.

Clients: name of socket @ hostname

**mode** Tells Sopen the type of socket being opened

**s** open a server

**S** open a server (but, on failure, retry)

**c** open a client

**s###** open a server with the specified port number

**c###** open a client to a server with the specified port number

**returns** Servers: a 3-tuple Socket\* - use Saccept() with it to get a full duplex communications socket

Clients: a 5-tuple Socket\*

# SSL: Saccept()

**Socket \*Saccept(Socket \*skt)**

**skt** a Socket\* server, opened by Sopen()

# SSL: Saccept()

**Socket \*Saccept(Socket \*skt)**

**skt** a Socket\* server, opened by Sopen()

**returns** a full 5-tuple Socket\* accept socket, useful for full duplex communications

# SSL: Sclose()

**void Sclose(Socket \*skt)**

Use this function to close a Sopen()ed Socket.



# SSL: Single Server Example

```
#include "sockets.h"
Socket *server;
Socket *skt;
server= Sopen("servername", "s");
skt    = Saccept(server);
...
Sclose(skt);
Sclose(server);
```

# SSL: Multiple Accepts Example

```
#include "sockets.h"
Socket *server= NULL;
Socket *skt    = NULL;
server= Sopen("servername", "s");
do {
    skt= Saccept(server);
    ...
    Sclose(skt);
} while(whatever);
Sclose(server);
```

# SSL: A Client Example

```
#include "sockets.h"
Socket *client;
while(1) {
    client= Sopen("servername", "c");
    if(client) break;
    sleep(1);
}
...
Sclose(client);
```

- If the “servername” server is on the local machine, just the name may be used
- For servers on non-local machines, use “servername@hostname”

# SSL: I/O

<i>Return Type</i>	<i>Function Name</i>	<i>Argument List</i>
char *	Sgets	(char *buf, int maxbuf, Socket *skt)
int	Speek	(Socket *skt, char *buf, int buflen)
void	Sprintf	(Socket *skt, char *fmt,...)
void	Sputs	(char *buf, Socket *skt)
int	Sread	(Socket *skt, char *buf, int buflen)
int	Sreadbytes	(Socket *skt, char *buf, int buflen)
int	Sscanf	(Socket *skt, char *fmt, ...)
int	Svprintf	(Socket *skt, char *fmt, void *args)

- Sread() reads up to buflen bytes, returning only whatever is currently present on the Socket

# SSL: I/O

<i>Return Type</i>	<i>Function Name</i>	<i>Argument List</i>
char *	Sgets	(char *buf, int maxbuf, Socket *skt)
int	Speek	(Socket *skt, char *buf, int buflen)
void	Sprintf	(Socket *skt, char *fmt,...)
void	Sputs	(char *buf, Socket *skt)
int	Sread	(Socket *skt, char *buf, int buflen)
int	Sreadbytes	(Socket *skt, char *buf, int buflen)
int	Sscanf	(Socket *skt, char *fmt, ...)
int	Svprintf	(Socket *skt, char *fmt, void *args)

- Sread() reads up to buflen bytes, returning only whatever is currently present on the Socket
- Sreadbytes() requires buflen bytes, and will block until it gets them.

(see ssl.c)

# Input/Output Multiplexing

How to handle multiple inputs:

1. polling : use `fcntl()` with `FNDELAY` or `ioctl` with `FIONBIO` to set the socket/file descriptor to non-blocking.

Repeatedly examine the descriptor for new data (ie. attempt to read).

Wastes cpu.

# Input/Output Multiplexing

How to handle multiple inputs:

1. polling : use `fcntl()` with `FNDELAY` or `ioctl` with `FIONBIO` to set the socket/file descriptor to non-blocking.

Repeatedly examine the descriptor for new data (ie. attempt to read).

Wastes cpu.

2. Use `fork()` and assign a separate process to handle each stream/descriptor.

Works with pipes, fifos, message queues, sockets, etc.

Use semaphores, shared memory, and signals to communicate between parent and child.

May implement callback functions.

# Input/Output Multiplexing

How to handle multiple inputs:

1. polling : use `fcntl()` with `FNDELAY` or `ioctl` with `FIONBIO` to set the socket/file descriptor to non-blocking.  
Repeatedly examine the descriptor for new data (ie. attempt to read).  
Wastes cpu.
2. Use `fork()` and assign a separate process to handle each stream/descriptor.  
Works with pipes, fifos, message queues, sockets, etc.  
Use semaphores, shared memory, and signals to communicate between parent and child.  
May implement callback functions.
3. Use threads; otherwise much like Method#2.



# Input/Output Multiplexing

How to handle multiple inputs:

1. polling : use `fcntl()` with `FNDELAY` or `ioctl` with `FIONBIO` to set the socket/file descriptor to non-blocking.  
Repeatedly examine the descriptor for new data (ie. attempt to read).  
Wastes cpu.
2. Use `fork()` and assign a separate process to handle each stream/descriptor.  
Works with pipes, fifos, message queues, sockets, etc.  
Use semaphores, shared memory, and signals to communicate between parent and child.  
May implement callback functions.
3. Use threads; otherwise much like Method#2.
4. Use `select()` (BSD) or `poll()` (SysV, streams only). This is a powerful method for multiplexing sockets and other file descriptors – call `select` or `poll` and put your process/thread to sleep until an event occurs (typically, receipt of data or a timeout).

# Multi I/O: select()

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);
```

```
int FD_ISSET(int fd, fd_set *set);
```

```
void FD_SET(int fd, fd_set *set);
```

```
void FD_ZERO(fd_set *set);
```

- select() works with “file descriptor sets”

Blocks until one or more of the file descriptors becomes “ready” for i/o

# Multi I/O: select()

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);
```

```
int FD_ISSET(int fd, fd_set *set);
```

```
void FD_SET(int fd, fd_set *set);
```

```
void FD_ZERO(fd_set *set);
```

- select() works with “file descriptor sets”

Blocks until one or more of the file descriptors becomes “ready” for i/o

- FD\_ZERO: clear all bits in the set

# Multi I/O: select()

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);
```

```
int FD_ISSET(int fd, fd_set *set);
```

```
void FD_SET(int fd, fd_set *set);
```

```
void FD_ZERO(fd_set *set);
```

- select() works with “file descriptor sets”  
Blocks until one or more of the file descriptors becomes “ready” for i/o
- FD\_ZERO: clear all bits in the set
- FD\_CLR: remove the specified file descriptor `fd` from the set (ie. off)

# Multi I/O: select()

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);
```

```
int FD_ISSET(int fd, fd_set *set);
```

```
void FD_SET(int fd, fd_set *set);
```

```
void FD_ZERO(fd_set *set);
```

- select() works with “file descriptor sets”  
Blocks until one or more of the file descriptors becomes “ready” for i/o
- FD\_ZERO: clear all bits in the set
- FD\_CLR: remove the specified file descriptor fd from the set (ie. off)
- FD\_SET: add the specified file descriptor fd to the set (ie. on)

# Multi I/O: select()

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);
```

```
int FD_ISSET(int fd, fd_set *set);
```

```
void FD_SET(int fd, fd_set *set);
```

```
void FD_ZERO(fd_set *set);
```

- select() works with “file descriptor sets”  
Blocks until one or more of the file descriptors becomes “ready” for i/o
- FD\_ZERO: clear all bits in the set
- FD\_CLR: remove the specified file descriptor `fd` from the set (ie. off)
- FD\_SET: add the specified file descriptor `fd` to the set (ie. on)
- FD\_ISSET: query if the bit associated with the specified file descriptor is set or not

# Multi I/O: select()

```
#include <sys/select.h>
```

```
int select(int nfd, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);
```

```
int FD_ISSET(int fd, fd_set *set);
```

```
void FD_SET(int fd, fd_set *set);
```

```
void FD_ZERO(fd_set *set);
```

- select() works with “file descriptor sets”  
Blocks until one or more of the file descriptors becomes “ready” for i/o
- FD\_ZERO: clear all bits in the set
- FD\_CLR: remove the specified file descriptor fd from the set (ie. off)
- FD\_SET: add the specified file descriptor fd to the set (ie. on)
- FD\_ISSET: query if the bit associated with the specified file descriptor is set or not
- fd\_set: supports 32 to 1024 bits arranged in an array of long integers (via unions), though few kernels will actually support that many open files.

# Multi I/O: select()

```
#include <sys/select.h>
```

```
int select(int nfd, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);
```

```
int FD_ISSET(int fd, fd_set *set);
```

```
void FD_SET(int fd, fd_set *set);
```

```
void FD_ZERO(fd_set *set);
```

- select() works with “file descriptor sets”

Blocks until one or more of the file descriptors becomes “ready” for i/o

- FD\_ZERO: clear all bits in the set
- FD\_CLR: remove the specified file descriptor `fd` from the set (ie. off)
- FD\_SET: add the specified file descriptor `fd` to the set (ie. on)
- FD\_ISSET: query if the bit associated with the specified file descriptor is set or not
- `fd_set`: supports 32 to 1024 bits arranged in an array of long integers (via unions), though few kernels will actually support that many open files.
- `nfd`: should be one more than the maximum file descriptor value over all the sets



# Multi I/O: select(), con't.

```
#include <sys/select.h>
```

```
int select(int nfd, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

- timeout: specifies a maximum time for blocking

```
struct timeval {  
    long tv_sec; // seconds  
    long tv_usec; // microseconds  
};
```

# Multi I/O: select(), con't.

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

- timeout: specifies a maximum time for blocking

```
struct timeval {
```

```
    long tv_sec; // seconds
```

```
    long tv_usec; // microseconds
```

```
};
```

- A timeout with zeros for tv\_sec and tv\_usec set returns immediately

# Multi I/O: select(), con't.

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

- timeout: specifies a maximum time for blocking

```
struct timeval {
```

```
    long tv_sec; // seconds
```

```
    long tv_usec; // microseconds
```

```
};
```

- A timeout with zeros for tv\_sec and tv\_usec set returns immediately
- select() will update the timeout argument to indicate how much time was left

# Multi I/O: select(), con't.

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

- timeout: specifies a maximum time for blocking

```
struct timeval {
```

```
    long tv_sec; // seconds
```

```
    long tv_usec; // microseconds
```

```
};
```

- A timeout with zeros for tv\_sec and tv\_usec set returns immediately
- select() will update the timeout argument to indicate how much time was left
- readfds: determine if the file descriptors in this set are ready for reading (ie. a read will not block)

# Multi I/O: select(), con't.

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

- timeout: specifies a maximum time for blocking

```
struct timeval {
```

```
    long tv_sec; // seconds
```

```
    long tv_usec; // microseconds
```

```
};
```

- A timeout with zeros for tv\_sec and tv\_usec set returns immediately
- select() will update the timeout argument to indicate how much time was left
- readfds: determine if the file descriptors in this set are ready for reading (ie. a read will not block)
- writefds: determine if the file descriptors in this set are ready for writing (ie. a write will not block)

# Multi I/O: select(), con't.

```
#include <sys/select.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

- timeout: specifies a maximum time for blocking

```
struct timeval {
```

```
    long tv_sec; // seconds
```

```
    long tv_usec; // microseconds
```

```
};
```

- A timeout with zeros for tv\_sec and tv\_usec set returns immediately
- select() will update the timeout argument to indicate how much time was left
- readfds: determine if the file descriptors in this set are ready for reading (ie. a read will not block)
- writefds: determine if the file descriptors in this set are ready for writing (ie. a write will not block)
- exceptfds: determine if the file descriptors in this set have any exceptions (out-of-band)

# Multi I/O: select(), con't.

- A null timeout means block forever (until a descriptor is ready for i/o)

# Multi I/O: `select()`, con't.

- A null timeout means block forever (until a descriptor is ready for i/o)
- Note that signals will cause `select()` to return `-1/EINTR`



# Multi I/O: `select()`, con't.

- A null timeout means block forever (until a descriptor is ready for i/o)
- Note that signals will cause `select()` to return `-1/EINTR`
- on exit, the sets are modified in place to indicate which file descriptors are ready for i/o

# Multi I/O: `select()`, con't.

- A null timeout means block forever (until a descriptor is ready for i/o)
- Note that signals will cause `select()` to return -1/EINTR
- on exit, the sets are modified in place to indicate which file descriptors are ready for i/o
- `select()` returns
  - + the quantity of ready descriptors,
  - 0 on timeout,
  - 1 on failure (with `errno` suitably set)

(see *select.c*)

# Multi I/O: pselect()

```
int pselect(int nfd, fd_set *readfds, fd_set *writefds,  
            fd_set *exceptfds, const struct timespec *timeout,  
            const sigset_t *sigmask);
```

- The pselect() function resembles select()

# Multi I/O: pselect()

```
int pselect(int nfd, fd_set *readfds, fd_set *writefds,  
fd_set *exceptfds, const struct timespec *timeout,  
const sigset_t *sigmask);
```

- The pselect() function resembles select()
- sigmask: a pointer to a signal mask

# Multi I/O: pselect()

```
int pselect(int nfd, fd_set *readfds, fd_set *writefds,  
fd_set *exceptfds, const struct timespec *timeout,  
const sigset_t *sigmask);
```

- The `pselect()` function resembles `select()`
- `sigmask`: a pointer to a signal mask
- `pselect()` first replaces the current signal mask with `sigmask`, then does the `select()` function, then restores the original signal mask.

# Multi I/O: pselect()

```
int pselect(int nfd, fd_set *readfds, fd_set *writefds,  
fd_set *exceptfds, const struct timespec *timeout,  
const sigset_t *sigmask);
```

- The pselect() function resembles select()
- sigmask: a pointer to a signal mask
- pselect() first replaces the current signal mask with sigmask, then does the select() function, then restores the original signal mask.
- Equivalent to *atomically* executing the following: *(sigprocmask() is used to fetch and/or change the signal mask of the calling thread)*

```
sigprocmask(SIG_SETMASK, &sigmask, &origmask);
```

```
ready = select(nfd, &readfds, &writefds, &exceptfds, timeout);
```

```
sigprocmask(SIG_SETMASK, &origmask, NULL);
```

# Multi I/O: pselect()

```
int pselect(int nfd, fd_set *readfds, fd_set *writefds,  
fd_set *exceptfds, const struct timespec *timeout,  
const sigset_t *sigmask);
```

- The pselect() function resembles select()
- sigmask: a pointer to a signal mask
- pselect() first replaces the current signal mask with sigmask, then does the select() function, then restores the original signal mask.
- Equivalent to *atomically* executing the following: *(sigprocmask() is used to fetch and/or change the signal mask of the calling thread)*

```
sigprocmask(SIG_SETMASK, &sigmask, &origmask);
```

```
ready = select(nfd, &readfds, &writefds, &exceptfds, timeout);
```

```
sigprocmask(SIG_SETMASK, &origmask, NULL);
```

- This function is needed when one wants to handle both signals and file descriptors while preventing a race condition.

# Multi I/O: pselect()

```
int pselect(int nfd, fd_set *readfds, fd_set *writefds,  
fd_set *exceptfds, const struct timespec *timeout,  
const sigset_t *sigmask);
```

- The pselect() function resembles select()
- sigmask: a pointer to a signal mask
- pselect() first replaces the current signal mask with sigmask, then does the select() function, then restores the original signal mask.
- Equivalent to *atomically* executing the following: *(sigprocmask() is used to fetch and/or change the signal mask of the calling thread)*

```
sigprocmask(SIG_SETMASK, &sigmask, &origmask);  
ready = select(nfd, &readfds, &writefds, &exceptfds, timeout);  
sigprocmask(SIG_SETMASK, &origmask, NULL);
```

- This function is needed when one wants to handle both signals and file descriptors while preventing a race condition.
- struct timespec supports seconds and nanoseconds

```
struct timespec {  
  
    long tv_sec; // seconds  
  
    long tv_nsec; // nanoseconds
```



# Multi I/O: poll()

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfd, int timeout);
```

- Functionality is similar to that of select()

# Multi I/O: poll()

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfd, int timeout);
```

- Functionality is similar to that of select()
- fds the set of nfd file descriptors to be monitored:

```
struct pollfd {
```

```
    int fd; // file descriptors to be checked
```

```
    short events; // (bitwise) events of interest on fd
```

```
    short revents; // (bitwise) events that occurred on fd
```

```
};
```

# Multi I/O: poll()

**#include <poll.h>**

**int poll(struct pollfd \*fds, nfds\_t nfds, int timeout);**

- Functionality is similar to that of select()
- fds the set of nfds file descriptors to be monitored:

```
struct pollfd {
```

```
    int fd; // file descriptors to be checked
```

```
    short events; // (bitwise) events of interest on fd
```

```
    short revents; // (bitwise) events that occurred on fd
```

```
};
```

- Events

<i>Event</i>	<i>Input to</i>	<i>Result from</i>	<i>Description</i>
<i>Name</i>	<i>Events</i>	<i>Events</i>	<i>Description</i>
POLLIN	.	.	nonpriority msg present

# Multi I/O: poll()

**#include <poll.h>**

**int poll(struct pollfd \*fds, nfds\_t nfds, int timeout);**

- Functionality is similar to that of select()
- fds the set of nfds file descriptors to be monitored:

```
struct pollfd {
```

```
    int fd; // file descriptors to be checked
```

```
    short events; // (bitwise) events of interest on fd
```

```
    short revents; // (bitwise) events that occurred on fd
```

```
};
```

- Events

<i>Event Name</i>	<i>Input to Events</i>	<i>Result from Events</i>	<i>Description</i>
POLLIN	.	.	nonpriority msg present
POLLPRI	.	.	priority msg present

# Multi I/O: poll()

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);
```

- Functionality is similar to that of select()
- fds the set of nfds file descriptors to be monitored:

```
struct pollfd {
```

```
    int fd; // file descriptors to be checked
```

```
    short events; // (bitwise) events of interest on fd
```

```
    short revents; // (bitwise) events that occurred on fd
```

```
};
```

- Events

<i>Event</i> <i>Name</i>	<i>Input to</i> <i>Events</i>	<i>Result from</i> <i>Events</i>	<i>Description</i> <i>Description</i>
POLLIN	.	.	nonpriority msg present
POLLPRI	.	.	priority msg present
POLLOUT	.	.	fd is writable

# Multi I/O: poll()

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);
```

- Functionality is similar to that of select()
- fds the set of nfds file descriptors to be monitored:

```
struct pollfd {
```

```
    int fd; // file descriptors to be checked
```

```
    short events; // (bitwise) events of interest on fd
```

```
    short revents; // (bitwise) events that occurred on fd
```

```
};
```

- Events

<i>Event Name</i>	<i>Input to Events</i>	<i>Result from Events</i>	<i>Description</i>
POLLIN	.	.	nonpriority msg present
POLLPRI	.	.	priority msg present
POLLOUT	.	.	fd is writable
POLLERR		.	error msg arrived

# Multi I/O: poll()

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);
```

- Functionality is similar to that of select()
- fds the set of nfds file descriptors to be monitored:

```
struct pollfd {  
    int fd; // file descriptors to be checked  
    short events; // (bitwise) events of interest on fd  
    short revents; // (bitwise) events that occurred on fd  
};
```

- Events

<i>Event</i> <i>Name</i>	<i>Input to</i> <i>Events</i>	<i>Result from</i> <i>Events</i>	<i>Description</i> <i>Description</i>
POLLIN	.	.	nonpriority msg present
POLLPRI	.	.	priority msg present
POLLOUT	.	.	fd is writable
POLLERR		.	error msg arrived
POLLHUP		.	hangup has occurred

# Multi I/O: poll()

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);
```

- Functionality is similar to that of select()
- fds the set of nfds file descriptors to be monitored:

```
struct pollfd {  
    int fd; // file descriptors to be checked  
    short events; // (bitwise) events of interest on fd  
    short revents; // (bitwise) events that occurred on fd  
};
```

- Events

<i>Event Name</i>	<i>Input to Events</i>	<i>Result from Events</i>	<i>Description</i>
POLLIN	.	.	nonpriority msg present
POLLPRI	.	.	priority msg present
POLLOUT	.	.	fd is writable
POLLERR		.	error msg arrived
POLLHUP		.	hangup has occurred
POLLNVAL		.	fd not an open stream



# Multi I/O: poll(), con't.

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfd, int timeout);
```

- timeout: - → block forever (until an event)

# Multi I/O: poll(), con't.

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfd, int timeout);
```

- timeout: - → block forever (until an event)
- timeout: 0 → return immediately

# Multi I/O: poll(), con't.

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfd, int timeout);
```

- timeout: - → block forever (until an event)
- timeout: 0 → return immediately
- timeout: + → wait this many milliseconds

# Multi I/O: poll(), con't.

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfd, int timeout);
```

- timeout: - → block forever (until an event)
- timeout: 0 → return immediately
- timeout: + → wait this many milliseconds
- POLLIN: data ready to read

# Multi I/O: poll(), con't.

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfd, int timeout);
```

- timeout: - → block forever (until an event)
- timeout: 0 → return immediately
- timeout: + → wait this many milliseconds
- POLLIN: data ready to read
- POLLPRI: urgent data ready to read (ie. out-of-band data on a tcp socket)

# Multi I/O: poll(), con't.

```
#include <poll.h>
```

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);
```

- timeout: - → block forever (until an event)
- timeout: 0 → return immediately
- timeout: + → wait this many milliseconds
- POLLIN: data ready to read
- POLLPRI: urgent data ready to read (ie. out-of-band data on a tcp socket)
- Returns
  - 1 on failure,
  - + qty descriptors with non-zero revents field,
  - 0 for timeout, no descriptors ready

(see poll.c)

# Daemons

- A **daemon** is a process which is expected to execute for a long time.

# Daemons

- A **daemon** is a process which is expected to execute for a long time.
- Often started when the system is brought up by `init`



# Daemons

- A **daemon** is a process which is expected to execute for a long time.
- Often started when the system is brought up by init
- They have no controlling terminal. *(thus they aren't terminated by users logging in and out)*

# Daemons

- A **daemon** is a process which is expected to execute for a long time.
- Often started when the system is brought up by init
- They have no controlling terminal. *(thus they aren't terminated by users logging in and out)*
- The Simple Sockets Library's PortMaster is an example of a daemon process *(that is not started by init)*

# Daemons

- A **daemon** is a process which is expected to execute for a long time.
- Often started when the system is brought up by `init`
- They have no controlling terminal. *(thus they aren't terminated by users logging in and out)*
- The Simple Socket Library's PortMaster is an example of a daemon process *(that is not started by init)*
- Use BSD: `ps -ajx` or SYSV: `ps -efjc` to see a list of such processes on your system.

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.



# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.
3. The nascent daemon then calls `setsid()` to create a new “session”

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.
3. The nascent daemon then calls `setsid()` to create a new “session”
  - daemon becomes a session leader of a new session

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.
3. The nascent daemon then calls `setsid()` to create a new “session”
  - daemon becomes a session leader of a new session
  - daemon becomes a process group leader of a new process group

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.
3. The nascent daemon then calls `setsid()` to create a new “session”
  - daemon becomes a session leader of a new session
  - daemon becomes a process group leader of a new process group
  - daemon has no controlling terminal

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.
3. The nascent daemon then calls `setsid()` to create a new “session”
  - daemon becomes a session leader of a new session
  - daemon becomes a process group leader of a new process group
  - daemon has no controlling terminal
4. Change directory to the root directory

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.
3. The nascent daemon then calls `setsid()` to create a new “session”
  - daemon becomes a session leader of a new session
  - daemon becomes a process group leader of a new process group
  - daemon has no controlling terminal
4. Change directory to the root directory
  - allows mounted filesystems to be removed

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.
3. The nascent daemon then calls `setsid()` to create a new “session”
  - daemon becomes a session leader of a new session
  - daemon becomes a process group leader of a new process group
  - daemon has no controlling terminal
4. Change directory to the root directory
  - allows mounted filesystems to be removed
  - alternative: have a dedicated directory where the daemon's file-related work will be done

# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.
3. The nascent daemon then calls `setsid()` to create a new “session”
  - daemon becomes a session leader of a new session
  - daemon becomes a process group leader of a new process group
  - daemon has no controlling terminal
4. Change directory to the root directory
  - allows mounted filesystems to be removed
  - alternative: have a dedicated directory where the daemon's file-related work will be done
5. Unneeded file descriptors the child inherited from the parent process should be closed



# Daemon Coding Rules

1. Set file mode creation mask to zero (`umask`)
  - The inherited mask may override what the daemon sets
  - The daemon is thereby in control of all its read/write/execute permissions.
2. The parent process should call `fork()` and `exit`. (*shells will then act as if the command is finished*)
  - The child gets its own process ID, has a parent, and so doesn't start as a process group leader.
3. The nascent daemon then calls `setsid()` to create a new “session”
  - daemon becomes a session leader of a new session
  - daemon becomes a process group leader of a new process group
  - daemon has no controlling terminal
4. Change directory to the root directory
  - allows mounted filesystems to be removed
  - alternative: have a dedicated directory where the daemon's file-related work will be done
5. Unneeded file descriptors the child inherited from the parent process should be closed
6. (*optional*) use file and record locking (`lockfile()`) to permit only one instance of the running daemon

# Get/Set Resource Limits

```
#include <sys/time.h>
```

```
#include <sys/resource.h>
```

```
int getrlimit(int resource, struct rlimit *rlim);
```

```
int setrlimit(int resource, const struct rlimit *rlim);
```

- These two functions allow one to set or get resource limits

# Get/Set Resource Limits

```
#include <sys/time.h>
```

```
#include <sys/resource.h>
```

```
int getrlimit(int resource, struct rlimit *rlim);
```

```
int setrlimit(int resource, const struct rlimit *rlim);
```

- These two functions allow one to set or get resource limits
- Each resource limit is given a soft and hard limit

# Get/Set Resource Limits

```
#include <sys/time.h>
```

```
#include <sys/resource.h>
```

```
int getrlimit(int resource, struct rlimit *rlim);
```

```
int setrlimit(int resource, const struct rlimit *rlim);
```

- These two functions allow one to set or get resource limits
- Each resource limit is given a soft and hard limit  
    **soft** kernel enforces this limit

# Get/Set Resource Limits

```
#include <sys/time.h>
```

```
#include <sys/resource.h>
```

```
int getrlimit(int resource, struct rlimit *rlim);
```

```
int setrlimit(int resource, const struct rlimit *rlim);
```

- These two functions allow one to set or get resource limits
- Each resource limit is given a soft and hard limit
  - soft** kernel enforces this limit
  - hard** a ceiling for the soft limit

# Get/Set Resource Limits

```
#include <sys/time.h>
```

```
#include <sys/resource.h>
```

```
int getrlimit(int resource, struct rlimit *rlim);
```

```
int setrlimit(int resource, const struct rlimit *rlim);
```

- These two functions allow one to set or get resource limits
- Each resource limit is given a soft and hard limit
  - soft** kernel enforces this limit
  - hard** a ceiling for the soft limit
- struct rlimit {
  - rlim\_t rlim\_cur; /\* Soft limit \*/
  - rlim\_t rlim\_max; /\* Hard limit (ceiling for rlim\_cur) \*/
- };

# Get/Set Resource Limits

```
#include <sys/time.h>
```

```
#include <sys/resource.h>
```

```
int getrlimit(int resource, struct rlimit *rlim);
```

```
int setrlimit(int resource, const struct rlimit *rlim);
```

- These two functions allow one to set or get resource limits
- Each resource limit is given a soft and hard limit
  - soft** kernel enforces this limit
  - hard** a ceiling for the soft limit
- struct rlimit {
  - rlim\_t rlim\_cur; /\* Soft limit \*/
  - rlim\_t rlim\_max; /\* Hard limit (ceiling for rlim\_cur) \*/
- };
- A sampling of resource limits is included on the next slide

# Get/Set Resource Limits

- A sample of resource limits: (*see man getrlimit*)

RLIMIT_AS	max size of process's virtual memory, in bytes
RLIMIT_CORE	max size of a core file (when 0, no core dumps created)
RLIMIT_CPU	cpu time limit (sec)
RLIMIT_DATA	max size of process' data segment
RLIMIT_FSIZE	max size of files a process may create
RLIMIT_LOCKS	limits qty of flock+fcntl locks
RLIMIT_MEMLOCK	max qty of memory (bytes) that may be locked into ram
RLIMIT_MSGQUEUE	max qty of bytes allocated for POSIX message queues
RLIMIT_NICE	ceiling for a process' nice value ( <i>priority</i> )
<b><i>RLIMIT_NOFILE</i></b>	<b><i>specifies max file descriptor + 1 that may be opened</i></b>



# Get/Set Resource Limits

- A sample of resource limits: (*see man getrlimit*)

RLIMIT_AS	max size of process's virtual memory, in bytes
RLIMIT_CORE	max size of a core file (when 0, no core dumps created)
RLIMIT_CPU	cpu time limit (sec)
RLIMIT_DATA	max size of process' data segment
RLIMIT_FSIZE	max size of files a process may create
RLIMIT_LOCKS	limits qty of flock+fcntl locks
RLIMIT_MEMLOCK	max qty of memory (bytes) that may be locked into ram
RLIMIT_MSGQUEUE	max qty of bytes allocated for POSIX message queues
RLIMIT_NICE	ceiling for a process' nice value ( <i>priority</i> )
<b><i>RLIMIT_NOFILE</i></b>	<b><i>specifies max file descriptor + 1 that may be opened</i></b>

- So, for daemons: use `getrlimit() + RLIMIT_NOFILE` to inquire about the maximum qty of file descriptors that the daemon process may have inherited and close them all.

# Daemon Example

```
void daemonize(const char *cmd) {
    int i, fd0, fd1, fd2;
    pid_t pid;
    struct rlimit rl;
    struct sigaction sa;

    umask(0);                                /* clear file creation mask */
    if( getrlimit(RLIMIT_NOFILE,&rl) < 0) /* handle error */
    if      ((pid= fork()) < 0)              /* handle error */
    else if(pid != 0) exit(0);              /* parent process */
    setsid();
    sa.sa_handler= SIG_IGN;                  /* don't allow controlling TTYs */
    sigemptyset(&sa.sa_mask);
    sa.sa_flags= 0;
    if( sigaction(SIGHUP,&sa,NULL) < 0) /* handle error */
    if((pid= fork()) < 0)                   /* handle error */
    else if(pid != 0) exit(0);              /* parent */
    if(chdir("/") < 0)                      /* handle error */
    /* close all open file descriptors */
    if(rl.rlim_max == RLIM_INFINITY) rl.rlim_max= 1024;
    for(i= 0; i < rl.rlim_max; ++i) close(i);
    fd0= open("/dev/null",O_RDWR); fd1= dup(0); fd2= dup(0);
    openlog(cmd,LOG_CONS,LOG_DAEMON);
    if(fd0 != 0 || fd1 != 1 || fd2 != 2) {
        syslog(LOG_ERR,"unexpected_file_descriptors_%d_%d_%d",fd0,fd1,fd2);
        exit(1);
    }
}
```

# Error Logging: `openlog`

- Daemons can't use `stderr` or `stdout`; *there's no controlling terminal!*
- Solution: use `syslog()` to write messages to a *logging* socket on `/dev/log`

**`#include "syslog.h"`**

**`void openlog(const char *ident,int option,int facility)`**

**`openlog`** opens a connection to the `syslogd` daemon.

# Error Logging: `openlog`

- Daemons can't use `stderr` or `stdout`; *there's no controlling terminal!*
- Solution: use `syslog()` to write messages to a *logging* socket on `/dev/log`

**`#include "syslog.h"`**

**`void openlog(const char *ident,int option,int facility)`**

**`openlog`** opens a connection to the `syslogd` daemon.

**`ident`** This string is prepended to every message (usually the name of the originating program).

# Error Logging: `openlog`

- Daemons can't use `stderr` or `stdout`; *there's no controlling terminal!*
- Solution: use `syslog()` to write messages to a *logging* socket on `/dev/log`

**#include "syslog.h"**

**void `openlog`(const char \*`ident`,int `option`,int `facility`)**

**`openlog`** opens a connection to the `syslogd` daemon.

**`ident`** This string is prepended to every message (usually the name of the originating program).

**`option`** Controls writing target, connection handling, extra information

# Error Logging: `openlog`

- Daemons can't use `stderr` or `stdout`; *there's no controlling terminal!*
- Solution: use `syslog()` to write messages to a *logging* socket on `/dev/log`

**#include "syslog.h"**

**void `openlog(const char *ident,int option,int facility)`**

**`openlog`** opens a connection to the `syslogd` daemon.

**ident** This string is prepended to every message (usually the name of the originating program).

**option** Controls writing target, connection handling, extra information

<code>LOG_CONS</code>	Write directly to the system console when unable to write to system logger
-----------------------	--

# Error Logging: `openlog`

- Daemons can't use `stderr` or `stdout`; *there's no controlling terminal!*
- Solution: use `syslog()` to write messages to a *logging* socket on `/dev/log`

**#include "syslog.h"**

**void `openlog(const char *ident, int option, int facility)`**

**`openlog`** opens a connection to the `syslogd` daemon.

**ident** This string is prepended to every message (usually the name of the originating program).

**option** Controls writing target, connection handling, extra information

<code>LOG_CONS</code>	Write directly to the system console when unable to write to system logger
<code>LOG_NDELAY</code>	Open connection immediately

# Error Logging: `openlog`

- Daemons can't use `stderr` or `stdout`; *there's no controlling terminal!*
- Solution: use `syslog()` to write messages to a *logging* socket on `/dev/log`

**#include "syslog.h"**

**void `openlog(const char *ident, int option, int facility)`**

**`openlog`** opens a connection to the `syslogd` daemon.

**ident** This string is prepended to every message (usually the name of the originating program).

**option** Controls writing target, connection handling, extra information

<code>LOG_CONS</code>	Write directly to the system console when unable to write to system logger
<code>LOG_NDELAY</code>	Open connection immediately
<code>LOG_NOWAIT</code>	Do not wait for child processes (safer for signal handling)



# Error Logging: `openlog`

- Daemons can't use `stderr` or `stdout`; *there's no controlling terminal!*
- Solution: use `syslog()` to write messages to a *logging* socket on `/dev/log`

**#include "syslog.h"**

**void `openlog(const char *ident, int option, int facility)`**

**`openlog`** opens a connection to the `syslogd` daemon.

**ident** This string is prepended to every message (usually the name of the originating program).

**option** Controls writing target, connection handling, extra information

<code>LOG_CONS</code>	Write directly to the system console when unable to write to system logger
<code>LOG_NDELAY</code>	Open connection immediately
<code>LOG_NOWAIT</code>	Do not wait for child processes (safer for signal handling)
<code>LOG_ODELAY</code>	Delay connecting to <code>syslogd</code> until first message is logged

# Error Logging: `openlog`

- Daemons can't use `stderr` or `stdout`; *there's no controlling terminal!*
- Solution: use `syslog()` to write messages to a *logging* socket on `/dev/log`

**#include "syslog.h"**

**void `openlog(const char *ident, int option, int facility)`**

**`openlog`** opens a connection to the `syslogd` daemon.

**ident** This string is prepended to every message (usually the name of the originating program).

**option** Controls writing target, connection handling, extra information

<code>LOG_CONS</code>	Write directly to the system console when unable to write to system logger
<code>LOG_NDELAY</code>	Open connection immediately
<code>LOG_NOWAIT</code>	Do not wait for child processes (safer for signal handling)
<code>LOG_ODELAY</code>	Delay connecting to <code>syslogd</code> until first message is logged
<code>LOG_PERROR</code>	Also write message to <code>stderr</code>

# Error Logging: `openlog`

- Daemons can't use `stderr` or `stdout`; *there's no controlling terminal!*
- Solution: use `syslog()` to write messages to a *logging* socket on `/dev/log`

**#include "syslog.h"**

**void `openlog(const char *ident,int option,int facility)`**

**`openlog`** opens a connection to the `syslogd` daemon.

**ident** This string is prepended to every message (usually the name of the originating program).

**option** Controls writing target, connection handling, extra information

<code>LOG_CONS</code>	Write directly to the system console when unable to write to system logger
<code>LOG_NDELAY</code>	Open connection immediately
<code>LOG_NOWAIT</code>	Do not wait for child processes (safer for signal handling)
<code>LOG_ODELAY</code>	Delay connecting to <code>syslogd</code> until first message is logged
<code>LOG_PERROR</code>	Also write message to <code>stderr</code>
<code>LOG_PID</code>	Log pid with each message

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG\_AUTH            authorization programs such as login, su, etc

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at
LOG_DAEMON	system daemons (inetd, routed, etc)



# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at
LOG_DAEMON	system daemons (inetd, routed, etc)
LOG_FTP	ftpd daemon

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at
LOG_DAEMON	system daemons (inetd, routed, etc)
LOG_FTP	ftpd daemon
LOG_KERN	kernel messages

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at
LOG_DAEMON	system daemons (inetd, routed, etc)
LOG_FTP	ftpd daemon
LOG_KERN	kernel messages
LOG_LPR	line printer daemons (lpd, lpc, etc)

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at
LOG_DAEMON	system daemons (inetd, routed, etc)
LOG_FTP	ftpd daemon
LOG_KERN	kernel messages
LOG_LPR	line printer daemons (lpd, lpc, etc)
LOG_MAIL	mail system

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at
LOG_DAEMON	system daemons (inetd, routed, etc)
LOG_FTP	ftpd daemon
LOG_KERN	kernel messages
LOG_LPR	line printer daemons (lpd, lpc, etc)
LOG_MAIL	mail system
LOG_NEWS	usenet network news system

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at
LOG_DAEMON	system daemons (inetd, routed, etc)
LOG_FTP	ftpd daemon
LOG_KERN	kernel messages
LOG_LPR	line printer daemons (lpd, lpc, etc)
LOG_MAIL	mail system
LOG_NEWS	usenet network news system
LOG_SYSLOG	messages from syslogd itself

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at
LOG_DAEMON	system daemons (inetd, routed, etc)
LOG_FTP	ftpd daemon
LOG_KERN	kernel messages
LOG_LPR	line printer daemons (lpd, lpc, etc)
LOG_MAIL	mail system
LOG_NEWS	usenet network news system
LOG_SYSLOG	messages from syslogd itself
LOG_USER	messages from user processes ( <i>default</i> )

# Error Logging: openlog, con't.

**#include "syslog.h"**

**void openlog(const char \*ident,int option,int facility)**

**openlog** opens a connection to the syslogd daemon (*continued*)

**facility** Lets the configuration specify how messages from different facilities are to be handled. The following is a sample, not a complete list:

LOG_AUTH	authorization programs such as login, su, etc
LOG_CRON	cron and at
LOG_DAEMON	system daemons (inetd, routed, etc)
LOG_FTP	ftpd daemon
LOG_KERN	kernel messages
LOG_LPR	line printer daemons (lpd, lpc, etc)
LOG_MAIL	mail system
LOG_NEWS	usenet network news system
LOG_SYSLOG	messages from syslogd itself
LOG_USER	messages from user processes ( <i>default</i> )
LOG_UUCP	uucp system



# Error Logging: syslog

**#include "syslog.h"**

**void syslog(int priority, const char \*format, ...)**

**syslog** Call this to generate a log message

**priority** Priorities; setlogmask() may be used to restrict logging to specified levels of priority

# Error Logging: syslog

**#include "syslog.h"**

**void syslog(int priority, const char \*format, ...)**

**syslog** Call this to generate a log message

**priority** Priorities; setlogmask() may be used to restrict logging to specified levels of priority

LOG\_EMERG      emergency (*system unusable*) (*highest priority*)

# Error Logging: syslog

**#include "syslog.h"**

**void syslog(int priority, const char \*format, ...)**

**syslog** Call this to generate a log message

**priority** Priorities; setlogmask() may be used to restrict logging to specified levels of priority

LOG\_EMERG      emergency (*system unusable*) (*highest priority*)

LOG\_ALERT      condition must be fixed immediately

# Error Logging: syslog

**#include "syslog.h"**

**void syslog(int priority, const char \*format, ...)**

**syslog** Call this to generate a log message

**priority** Priorities; setlogmask() may be used to restrict logging to specified levels of priority

LOG_EMERG	emergency ( <i>system unusable</i> ) ( <i>highest priority</i> )
LOG_ALERT	condition must be fixed immediately
LOG_CRIT	critical condition ( <i>ie. hard device error</i> )

# Error Logging: syslog

**#include "syslog.h"**

**void syslog(int priority, const char \*format, ...)**

**syslog** Call this to generate a log message

**priority** Priorities; setlogmask() may be used to restrict logging to specified levels of priority

LOG_EMERG	emergency ( <i>system unusable</i> ) ( <i>highest priority</i> )
LOG_ALERT	condition must be fixed immediately
LOG_CRIT	critical condition ( <i>ie. hard device error</i> )
LOG_ERR	error condition

# Error Logging: syslog

**#include “syslog.h”**

**void syslog(int priority,const char \*format,...)**

**syslog** Call this to generate a log message

**priority** Priorities; setlogmask() may be used to restrict logging to specified levels of priority

LOG_EMERG	emergency ( <i>system unusable</i> ) ( <i>highest priority</i> )
LOG_ALERT	condition must be fixed immediately
LOG_CRIT	critical condition ( <i>ie. hard device error</i> )
LOG_ERR	error condition
LOG_WARNING	warning condition

# Error Logging: syslog

**#include “syslog.h”**

**void syslog(int priority,const char \*format,...)**

**syslog** Call this to generate a log message

**priority** Priorities; setlogmask() may be used to restrict logging to specified levels of priority

LOG_EMERG	emergency ( <i>system unusable</i> ) ( <i>highest priority</i> )
LOG_ALERT	condition must be fixed immediately
LOG_CRIT	critical condition ( <i>ie. hard device error</i> )
LOG_ERR	error condition
LOG_WARNING	warning condition
LOG_NOTICE	normal <i>but significant</i> condition

# Error Logging: syslog

**#include "syslog.h"**

**void syslog(int priority, const char \*format, ...)**

**syslog** Call this to generate a log message

**priority** Priorities; setlogmask() may be used to restrict logging to specified levels of priority

LOG_EMERG	emergency ( <i>system unusable</i> ) ( <i>highest priority</i> )
LOG_ALERT	condition must be fixed immediately
LOG_CRIT	critical condition ( <i>ie. hard device error</i> )
LOG_ERR	error condition
LOG_WARNING	warning condition
LOG_NOTICE	normal <i>but significant</i> condition
LOG_INFO	informational



# Error Logging: syslog

**#include "syslog.h"**

**void syslog(int priority, const char \*format, ...)**

**syslog** Call this to generate a log message

**priority** Priorities; setlogmask() may be used to restrict logging to specified levels of priority

LOG_EMERG	emergency ( <i>system unusable</i> ) ( <i>highest priority</i> )
LOG_ALERT	condition must be fixed immediately
LOG_CRIT	critical condition ( <i>ie. hard device error</i> )
LOG_ERR	error condition
LOG_WARNING	warning condition
LOG_NOTICE	normal <i>but significant</i> condition
LOG_INFO	informational
LOG_DEBUG	debug message ( <i>lowest priority</i> )

# Error Logging: setlogmask

```
#include "syslog.h"
```

```
int setlogmask(int maskpri)
```

**setlogmask** *priority* *Priorities*; setlogmask() may be used to restrict logging to specified levels of priority

- Processes possess a *log priority mask* that filters which messages syslog() will actually log

# Error Logging: setlogmask

```
#include "syslog.h"
```

```
int setlogmask(int maskpri)
```

**setlogmask** *priority* *Priorities*; setlogmask() may be used to restrict logging to specified levels of priority

- Processes possess a *log priority mask* that filters which messages syslog() will actually log
- Other calls to syslog() will be ignored.

# Error Logging: setlogmask

```
#include "syslog.h"
```

```
int setlogmask(int maskpri)
```

**setlogmask** *priority* *Priorities*; setlogmask() may be used to restrict logging to specified levels of priority

- Processes possess a *log priority mask* that filters which messages syslog() will actually log
- Other calls to syslog() will be ignored.
- The mask has a bit assigned to each priority.

# Error Logging: setlogmask

```
#include "syslog.h"
```

```
int setlogmask(int maskpri)
```

**setlogmask** *priority* *Priorities*; setlogmask() may be used to restrict logging to specified levels of priority

- Processes possess a *log priority mask* that filters which messages syslog() will actually log
- Other calls to syslog() will be ignored.
- The mask has a bit assigned to each priority.
- The function sets the new *maskpri* mask and returns the prior mask

# Error Logging: setlogmask

```
#include "syslog.h"
```

```
int setlogmask(int maskpri)
```

**setlogmask** *priority* *Priorities*; setlogmask() may be used to restrict logging to specified levels of priority

- Processes possess a *log priority mask* that filters which messages syslog() will actually log
- Other calls to syslog() will be ignored.
- The mask has a bit assigned to each priority.
- The function sets the new maskpri mask and returns the prior mask
- Set up the mask by bitwise or-ing your selection of LOG\_EMERG, . . . LOG\_DEBUG priorities.

# Error Logging: setlogmask

```
#include "syslog.h"
```

```
int setlogmask(int maskpri)
```

**setlogmask** *priority* *Priorities*; setlogmask() may be used to restrict logging to specified levels of priority

- Processes possess a *log priority mask* that filters which messages syslog() will actually log
- Other calls to syslog() will be ignored.
- The mask has a bit assigned to each priority.
- The function sets the new maskpri mask and returns the prior mask
- Set up the mask by bitwise or-ing your selection of LOG\_EMERG, . . . LOG\_DEBUG priorities.
- Setting the mask to 0 will have no effect. (*other than returning the current logging priority mask*)

# Error Logging: **closelog**

**#include "syslog.h"**

**void closelog(void)**

**closelog** closes the descriptor being used to write to syslog. *(using this is optional)*



# Error Logging: Conventions

If the daemon uses a lock file, its generally located in `/var/run/name.pid`, where `name` is the daemon's name

# Error Logging: Conventions

If the daemon uses a lock file, its generally located in `/var/run/name.pid`, where `name` is the daemon's name

Daemon configuration options are usually stored in `/etc/name.conf`.

# Error Logging: Conventions

If the daemon uses a lock file, its generally located in `/var/run/name.pid`, where `name` is the daemon's name

Daemon configuration options are usually stored in `/etc/name.conf`.

Daemons are usually started from one of the system initialization scripts (`/etc/rc*` *or* `/etc/init.d/*`)

# Error Logging: Conventions

If the daemon uses a lock file, its generally located in `/var/run/name.pid`, where `name` is the daemon's name

Daemon configuration options are usually stored in `/etc/name.conf`.

Daemons are usually started from one of the system initialization scripts (`/etc/rc*` *or* `/etc/init.d/*`)

Daemons that need to restart automatically upon exiting can be restarted by including a `respawn` entry for it in `/etc/inittab/name.conf`.

# Error Logging: Conventions

If the daemon uses a lock file, its generally located in `/var/run/name.pid`, where `name` is the daemon's name

Daemon configuration options are usually stored in `/etc/name.conf`.

Daemons are usually started from one of the system initialization scripts (`/etc/rc*` *or* `/etc/init.d/*`)

Daemons that need to restart automatically upon exiting can be restarted by including a `respawn` entry for it in `/etc/inittab/name.conf`.

Daemons typically use the `SIGHUP` signal to tell them to re-read their configuration file. (*as daemons don't have terminals this signal is safe them to use in this fashion*)

# Unix Initialization

- Use `who -r` or `runlevel` from the command line to determine the current runlevel.

# Unix Initialization

- Use `who -r` or `runlevel` from the command line to determine the current runlevel.
- Use `init 1` to change to runlevel 1 (*as the superuser*)

# Unix Initialization

- Use `who -r` or `runlevel` from the command line to determine the current runlevel.
- Use `init 1` to change to runlevel 1 (*as the superuser*)
- To change the default runlevel to 5, for example, edit `/etc/inittab` and include the line  
`id:5:initdefault:`

Save and close the file, then issue the `reboot` command



# Unix Initialization

- Use `who -r` or `runlevel` from the command line to determine the current runlevel.
- Use `init 1` to change to runlevel 1 (*as the superuser*)
- To change the default runlevel to 5, for example, edit `/etc/inittab` and include the line  
`id:5:initdefault:`  
Save and close the file, then issue the `reboot` command
- Some services cannot be used until the system has certain other services running; runlevels provide this organization.

# Unix Initialization

- Use `who -r` or `runlevel` from the command line to determine the current runlevel.
- Use `init 1` to change to runlevel 1 (*as the superuser*)
- To change the default runlevel to 5, for example, edit `/etc/inittab` and include the line  
`id:5:initdefault:`  
Save and close the file, then issue the `reboot` command
- Some services cannot be used until the system has certain other services running; runlevels provide this organization.
- Use `shutdown -h now` (*as superuser*) to bring the system down

# Unix Initialization

- Use `who -r` or `runlevel` from the command line to determine the current runlevel.
- Use `init 1` to change to runlevel 1 (*as the superuser*)
- To change the default runlevel to 5, for example, edit `/etc/inittab` and include the line  
`id:5:initdefault:`  
Save and close the file, then issue the `reboot` command
- Some services cannot be used until the system has certain other services running; runlevels provide this organization.
- Use `shutdown -h now` (*as superuser*) to bring the system down
- Use `shutdown -r now` (*as superuser*) to reboot the system

# Unix Init: Runlevels

**Runlevel** The first six runlevels appear when Linux is “coming up”.

0. Halt (*do not set initdefault to this!*)

# Unix Init: Runlevels

**Runlevel** The first six runlevels appear when Linux is “coming up”.

0. Halt (*do not set initdefault to this!*)
1. Single User Mode

# Unix Init: Runlevels

**Runlevel** The first six runlevels appear when Linux is “coming up”.

0. Halt (*do not set initdefault to this!*)
1. Single User Mode
2. Multi-user (*without NFS, same as 3 without networking*)

# Unix Init: Runlevels

**Runlevel** The first six runlevels appear when Linux is “coming up”.

0. Halt (*do not set initdefault to this!*)
1. Single User Mode
2. Multi-user (*without NFS, same as 3 without networking*)
3. Full multi-user

# Unix Init: Runlevels

**Runlevel** The first six runlevels appear when Linux is “coming up”.

0. Halt (*do not set initdefault to this!*)
1. Single User Mode
2. Multi-user (*without NFS, same as 3 without networking*)
3. Full multi-user
4. unused



# Unix Init: Runlevels

**Runlevel** The first six runlevels appear when Linux is “coming up”.

0. Halt (*do not set initdefault to this!*)
1. Single User Mode
2. Multi-user (*without NFS, same as 3 without networking*)
3. Full multi-user
4. unused
5. X11

# Unix Init: Runlevels

**Runlevel** The first six runlevels appear when Linux is “coming up”.

0. Halt (*do not set initdefault to this!*)
1. Single User Mode
2. Multi-user (*without NFS, same as 3 without networking*)
3. Full multi-user
4. unused
5. X11
6. reboot (*do not set initdefault to this!*)

# Unit Init: Starting/Stopping Daemons

- After the Linux kernel has booted, the init program reads `/etc/inittab` (or `/etc/init/rcS.conf`)

# Unit Init: Starting/Stopping Daemons

- After the Linux kernel has booted, the init program reads `/etc/inittab` (or `/etc/init/rcS.conf`)
- Individual runlevels are then started by `/etc/init/rc.conf`

# Unit Init: Starting/Stopping Daemons

- After the Linux kernel has booted, the init program reads `/etc/inittab` (or `/etc/init/rcS.conf`)
- Individual runlevels are then started by `/etc/init/rc.conf`
- In `/etc/rc.d/rc#.d`, where `#` is the current runlevel, are scripts which will be executed in alphabetical order. Consequently, they have names such as `K46radvd`, `K50dnsmasq`, etc: the leading letter+number combination provides sequence control.

# Unit Init: Starting/Stopping Daemons

- After the Linux kernel has booted, the init program reads `/etc/inittab` (or `/etc/init/rcS.conf`)
- Individual runlevels are then started by `/etc/init/rc.conf`
- In `/etc/rc.d/rc#.d`, where `#` is the current runlevel, are scripts which will be executed in alphabetical order. Consequently, they have names such as `K46radvd`, `K50dnsmasq`, etc: the leading letter+number combination provides sequence control.
- Each of these `/etc/rc.d/rc#.d/*` scripts control the starting and stopping of a service (*system daemon*)

# Unit Init: Starting/Stopping Daemons

- After the Linux kernel has booted, the init program reads `/etc/inittab` (or `/etc/init/rcS.conf`)
- Individual runlevels are then started by `/etc/init/rc.conf`
- In `/etc/rc.d/rc#.d`, where `#` is the current runlevel, are scripts which will be executed in alphabetical order. Consequently, they have names such as `K46radvd`, `K50dnsmasq`, etc: the leading letter+number combination provides sequence control.
- Each of these `/etc/rc.d/rc#.d/*` scripts control the starting and stopping of a service (*system daemon*)
- Scripts beginning with “K” are *kill* scripts; those starting with “S” are *start* scripts

# Unit Init: Starting/Stopping Daemons

- After the Linux kernel has booted, the init program reads `/etc/inittab` (or `/etc/init/rcS.conf`)
- Individual runlevels are then started by `/etc/init/rc.conf`
- In `/etc/rc.d/rc#.d`, where `#` is the current runlevel, are scripts which will be executed in alphabetical order. Consequently, they have names such as `K46radvd`, `K50dnsmasq`, etc: the leading letter+number combination provides sequence control.
- Each of these `/etc/rc.d/rc#.d/*` scripts control the starting and stopping of a service (*system daemon*)
- Scripts beginning with “K” are *kill* scripts; those starting with “S” are *start* scripts
- When the system begins a new runlevel, all the files beginning with “K” are executed, passing the “stop” parameter to the associated daemon



# Unit Init: Starting/Stopping Daemons

- After the Linux kernel has booted, the init program reads `/etc/inittab` (or `/etc/init/rcS.conf`)
- Individual runlevels are then started by `/etc/init/rc.conf`
- In `/etc/rc.d/rc#.d`, where `#` is the current runlevel, are scripts which will be executed in alphabetical order. Consequently, they have names such as `K46radvd`, `K50dnsmasq`, etc: the leading letter+number combination provides sequence control.
- Each of these `/etc/rc.d/rc#.d/*` scripts control the starting and stopping of a service (*system daemon*)
- Scripts beginning with “K” are *kill* scripts; those starting with “S” are *start* scripts
- When the system begins a new runlevel, all the files beginning with “K” are executed, passing the “stop” parameter to the associated daemon
- When the system begins a new runlevel, all the files beginning with “S” are then executed, passing the “start” parameter to the associated daemon

# Unit Init: Starting/Stopping Daemons

- After the Linux kernel has booted, the init program reads `/etc/inittab` (or `/etc/init/rcS.conf`)
- Individual runlevels are then started by `/etc/init/rc.conf`
- In `/etc/rc.d/rc#.d`, where `#` is the current runlevel, are scripts which will be executed in alphabetical order. Consequently, they have names such as `K46radvd`, `K50dnsmasq`, etc: the leading letter+number combination provides sequence control.
- Each of these `/etc/rc.d/rc#.d/*` scripts control the starting and stopping of a service (*system daemon*)
- Scripts beginning with “K” are *kill* scripts; those starting with “S” are *start* scripts
- When the system begins a new runlevel, all the files beginning with “K” are executed, passing the “stop” parameter to the associated daemon
- When the system begins a new runlevel, all the files beginning with “S” are then executed, passing the “start” parameter to the associated daemon
- When one leaves a runlevel, nothing happens – only when a runlevel is started